

# Introduction to time series

ANOMALY DETECTION IN PYTHON



**Bekhruz (Bex) Tuychiev**

Kaggle Master, Data Science Content  
Creator

# Google stocks dataset

```
import pandas as pd
```

```
google = pd.read_csv("google.csv")  
google.head()
```

	Open	High	Low	Close	Volume
Date					
2006-01-03	211.47	218.05	209.32	217.83	13137450
2006-01-04	222.17	224.70	220.09	222.84	15292353
2006-01-05	223.22	226.00	220.97	225.85	10815661
2006-01-06	228.66	235.49	226.85	233.06	17759521
2006-01-09	233.44	236.94	230.70	233.68	12795837

# DateTime datatype

```
print(google['Date'].dtype)
```

```
object
```

```
google['Date'] = pd.to_datetime(google['Date'])
```

```
print(google.dtypes)
```

```
Date      datetime64[ns]  
Open              float64  
High              float64  
...
```

# Extracting features

```
google['day_of_week'] = google['Date'].dt.day_of_week
google['day_of_month'] = google['Date'].dt.day
google['month'] = google['Date'].dt.month

google.sample(5)
```

Date	Open	Low	Close	Volume	day_of_week	month	day_of_month
2016-02-29	721.00	716.84	717.22	2237474	0	2	29
2007-01-24	242.46	241.89	249.78	6074077	2	1	24
2007-05-24	237.81	235.99	237.40	4200474	3	5	24
2008-09-16	213.19	212.96	221.69	6991767	1	9	16
2008-03-31	218.04	216.22	220.46	4446368	0	3	31

# DatetimeIndex

```
google.set_index("Date", inplace=True)
```

```
google.head()
```

	Open	High	Low	Close	Volume	day_of_week	month	\
Date								
2006-01-03	211.47	218.05	209.32	217.83	13137450	1	1	
2006-01-04	222.17	224.70	220.09	222.84	15292353	2	1	
2006-01-05	223.22	226.00	220.97	225.85	10815661	3	1	
2006-01-06	228.66	235.49	226.85	233.06	17759521	4	1	
2006-01-09	233.44	236.94	230.70	233.68	12795837	0	1	

# Choosing periods

```
google["2008": "2010"].head()
```

	Open	High	Low	Close	Volume	day_of_week	month	\
Date								
2008-01-02	346.78	349.03	339.20	342.94	4306848	2	1	
2008-01-03	342.97	343.77	338.60	343.01	3252846	3	1	
2008-01-04	340.18	340.82	327.83	328.83	5359834	4	1	
2008-01-07	327.30	331.47	318.99	324.95	6404945	0	1	
2008-01-08	326.83	330.31	315.82	316.16	5341949	1	1	

# Choosing periods

```
google["2012-03": "2015-10-04"].head()
```

	Open	High	Low	Close	Volume	day_of_week	month	\
Date								
2012-03-01	311.44	313.16	309.38	311.51	2238010	3	3	
2012-03-02	311.31	312.31	310.47	310.94	1573214	4	3	
2012-03-05	310.53	311.56	306.00	307.43	1593250	0	3	
2012-03-06	304.33	304.71	297.22	302.78	3175216	1	3	
2012-03-07	304.83	305.90	303.23	303.70	1264892	2	3	

# Loading datasets with a DatetimeIndex

```
google = pd.read_csv("google.csv", parse_dates=["Date"], index_col="Date")

google.head()
```

	Open	High	Low	Close	Volume
Date					
2006-01-03	211.47	218.05	209.32	217.83	13137450
2006-01-04	222.17	224.70	220.09	222.84	15292353
2006-01-05	223.22	226.00	220.97	225.85	10815661
2006-01-06	228.66	235.49	226.85	233.06	17759521
2006-01-09	233.44	236.94	230.70	233.68	12795837



# Plotting time series

```
import matplotlib.pyplot as plt


google["Close"].plot(color="red")

plt.title("""Closing prices of Google
stocks from 2006 to 2018.""")
plt.show()
```




# Plotting time series

```
google["2010": "2010-07-01"]['Close'].plot(color="green")  
  
plt.title("Closing prices of Google stocks from January 2010 to July 2010.")  
plt.show()
```



# Plotting time series

```
google['Volume'].plot(color='red', figsize=(12, 4))  
  
plt.title("The number of traded Google stocks from 2006 to 2018.")
```



# MAD on time series

```
from pyod.models.mad import MAD

mad = MAD().fit(google[['Volume']])

is_outlier = mad.labels_ == 1

print(len(google[is_outlier]))
```

236

# IForest on time series

```
google['day_of_week'] = google.index.day_of_week
google['month'] = google.index.month
google['day_of_month'] = google.index.day

google.head()
```

	Open	High	Low	Close	Volume	day_of_week	month	\
Date								
2006-01-03	211.47	218.05	209.32	217.83	13137450	1	1	
2006-01-04	222.17	224.70	220.09	222.84	15292353	2	1	
2006-01-05	223.22	226.00	220.97	225.85	10815661	3	1	
2006-01-06	228.66	235.49	226.85	233.06	17759521	4	1	
2006-01-09	233.44	236.94	230.70	233.68	12795837	0	1	

# IForest on time series

```
from pyod.models.iforest import IForest

iforest = IForest().fit(google)

# Generate probabilities
probs = iforest.predict_proba(google)

# Isolate the outliers
is_outlier = probs[:, 1] > 0.75
outliers = google[is_outlier]

print(len(outliers))
```

# Let's practice!

ANOMALY DETECTION IN PYTHON

# Time Series Decomposition for Outlier Detection

ANOMALY DETECTION IN PYTHON



**Bekhruz (Bex) Tychiev**


Kaggle Master, Data Science Content  
Creator



# Seasonality

- Repeating patterns in the time series
- Has a fixed frequency:
  - hourly
  - daily
  - weekly
  - monthly, etc.
- Examples:
  - Daily temperatures
  - Ice-cream sales

# Seasonality



# seasonal\_decompose

```
from statsmodels.tsa.seasonal import seasonal_decompose


results = seasonal_decompose(google['Open'], period=365)

print(results)
```







```
<statsmodels.tsa.seasonal.DecomposeResult object at 0x7f0a67fac820>
```

# Plotting seasonality


```
results.seasonal.plot(color="red", figsize=(12, 4))
```



# Seasonality examples




# Initial plot of stocks




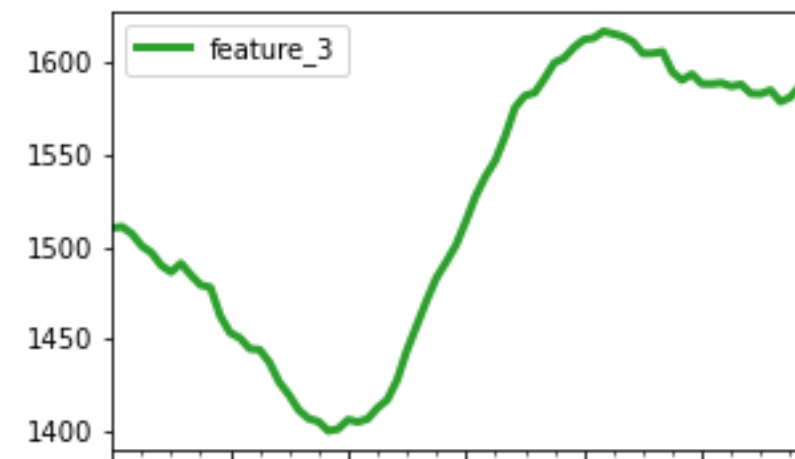

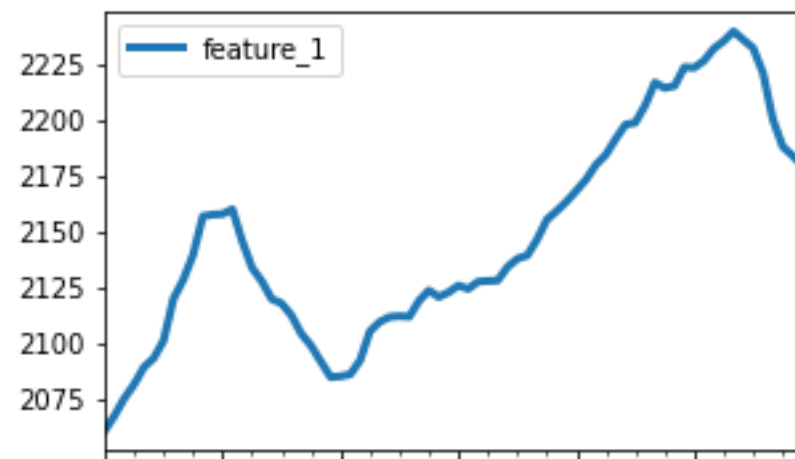


# Trend

```
results.trend.plot(color="red", figsize=(12, 4))
```




# Trend examples





# Residuals

```
results.resid.plot(color="red", figsize=(12, 4))
```




# Decomposition

```
figure = results.plot()
```

```
figure.set_figwidth(12)
```

```
figure.set_figheight(8)
```



# Fitting a classifier

```
# Extract and reshape residuals
results = seasonal_decompose(google['Volume'], period=365)
residuals = results.resid
residuals = residuals.values.reshape(-1, 1)

# Fit MAD
mad = MAD().fit(residuals)

# Find the outliers
is_outlier = mad.labels_ == 1
outliers = google[is_outlier]

print(len(outliers))
```

# Let's practice!

ANOMALY DETECTION IN PYTHON

# Outlier classifier ensembles

ANOMALY DETECTION IN PYTHON



**Bekhruz (Bex) Tychiev**

Kaggle Master, Data Science Content  
Creator

# Back to Airbnb

```
outliers = airbnb_df[iforest.labels_ == 1]

outlier_probs = iforest.predict_proba(outliers)
outlier_probs[:10]
```

```
array([[0.51999538, 0.48000462],
       [0.61789522, 0.38210478],
       [0.61802032, 0.38197968],
       [0.35184434, 0.64815566],
       [0.57533286, 0.42466714],
       [0.59038933, 0.40961067],
       [0.57677613, 0.42322387],
       [0.54158826, 0.45841174],
       [0.49118093, 0.50881907],
       [0.21387357, 0.78612643]])
```

# Probability threshold best practice

Threshold:

- 75% in low-risk cases
- >90% in high-cost cases like:
  - medicine
  - cyber security
  - limited data

# What is an ensemble?

- Combination of two or more classifiers
- Predictions are more stable



# Look at the data

```
google.head()
```

	Open	High	Low	Close	Volume	day_of_week	month	day
Date								
2006-01-03	211.47	218.05	209.32	217.83	13137450	1	1	3
2006-01-04	222.17	224.70	220.09	222.84	15292353	2	1	4
2006-01-05	223.22	226.00	220.97	225.85	10815661	3	1	5
2006-01-06	228.66	235.49	226.85	233.06	17759521	4	1	6
2006-01-09	233.44	236.94	230.70	233.68	12795837	0	1	9

# Scaling numeric features

```
from sklearn.preprocessing import QuantileTransformer

# Define the cols to be scaled
to_scale = ['Open', 'High', 'Low', 'Close', 'Volume']

# Initiate the transformer
qt = QuantileTransformer(output_distribution="normal")

# Scale and store the columns back
google.loc[:, to_scale] = qt.fit_transform(google[to_scale])
```

# Creating arrays

```
# Create a list of estimators
estimators = [KNN(n_neighbors=20), LOF(n_neighbors=20), IForest()]

# Create an empty array
shape = (len(google), len(estimators))
probability_scores = np.empty(shape=shape)
```

# Inside the loop

```
estimators = [KNN(n_neighbors=20), LOF(n_neighbors=20), IForest()]

shape = (len(google), len(estimators))
probability_scores = np.empty(shape=shape)

# Loop over and fit
for index, est in enumerate(estimators):
    est.fit(google)

    # Create probabilities
    probs = est.predict_proba(google)

    # Store the probs
    probability_scores[:, index] = probs[:, 1]
```

# Aggregating - mean

```
mean_scores = np.mean(probability_scores, axis=1)
```

```
mean_scores
```

```
array([0.20699869, 0.21455413, 0.17166271, ..., 0.31255075, 0.33553513,  
       0.32217186])
```

# Aggregating - median

```
median_scores = np.mean(probability_scores, axis=1)
```

```
median_scores
```

```
array([0.20699869, 0.21455413, 0.17166271, ..., 0.31255075, 0.33553513,  
       0.32217186])
```

# Probability filter

```
# Create a mask with 75% threshold
is_outlier = median_scores > 0.75

# Filter the outliers
outliers = google[is_outlier]

len(outliers)
```

3

# Summary of the steps

```
# Create a list of estimators
estimators = [KNN(n_neighbors=20), LOF(n_neighbors=20), IForest()]
probability_scores = np.empty(shape=(len(google), len(estimators)))

for index, est in enumerate(estimators):
    # Fit and generate probabilities
    est.fit(google)
    probs = est.predict_proba(google)

    # Store the probabilities
    probability_scores[:, index] = probs[:, 1]
```



# Summary of the steps

```
# Average the scores
mean_scores = np.mean(probability_scores, axis=1)

# Filter with 75% threshold
outliers = google[mean_scores > 0.75]

print(len(outliers))
```

3

# Let's practice!

ANOMALY DETECTION IN PYTHON

# How to deal with found outliers

ANOMALY DETECTION IN PYTHON



**Bekhruz (Bex) Tuychiev**

Kaggle Master, Data Science Content  
Creator

# Applications of anomaly detection

- Medicine
- Cyber security
- Fraud detection

Perform two analyses - with and without outliers.

# The reasons for outlier presence

- Data entry errors:
  - Typos
  - Measurement errors
  - Human mistakes
  - Drop unless fixed
- Sampling errors:
  - Not from the target distribution
  - Drop
- Natural:
  - Naturally odd but comes from the population
  - Do not drop

# Drop based on magnitude

- Too few: confirm and drop
- Too many: raises suspicion - use different models:
  - GLMs
  - Quantile Regression
  - GEEs
- Forms a cluster: perform deeper analysis

# Trimming

```
# Calculate the percentiles
percentile_first = google['Volume'].quantile(0.01)
percentile_99th = google['Volume'].quantile(0.99)

# Trim
google['Volume'] = google['Volume'].clip(percentile_first, percentile_99th)
```

# Replacing

```
google.replace(0, 100, inplace=True)
```



# Let's practice!

ANOMALY DETECTION IN PYTHON

# Congratulations!

ANOMALY DETECTION IN PYTHON



**Bekhruz (Bex) Tychiev**

Kaggle Master, Data Science Content  
Creator

# Chapters 1-2 recap

- Chapter 1 - univariate outlier detection:
  - Visual methods
  - Median Absolute Deviation of `pyod`
- Chapter 2 - Isolation Forest:
  - `IForest` of `pyod`
  - `iTrees`
  - Using outlier probabilities

# Chapters 3-4 recap

- Chapter 3 - distance and density-based algorithms:
  - `KNN` for outlier detection
  - `QuantileTransformer` for normalization
  - Local Outlier Factor algorithm
- Chapter 4 - time series anomalies and outlier ensembles:
  - Time series decomposition
  - Time series outlier detection from residuals
  - Building outlier ensembles manually
  - How to deal with found outliers

# Thank you!

ANOMALY DETECTION IN PYTHON