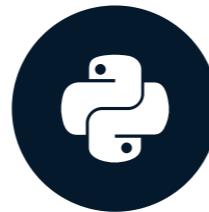


# From vectors to graphs

RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN



**Meri Nova**  
Machine Learning Engineer

# Vector RAG limitations

"How can I install Python locally?"

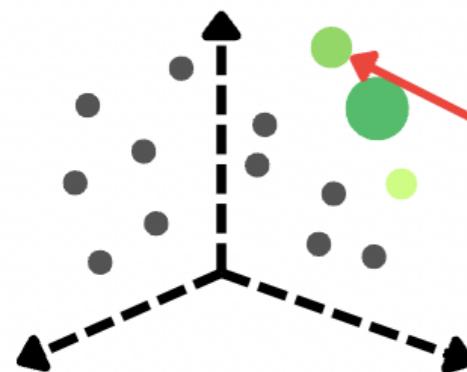
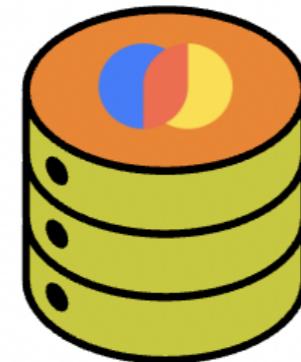


Embedding Model

[-0.829647...]



Vector Database



"Install Python in 3 easy steps?"



# Vector RAG limitations

"How can I install Python locally?"

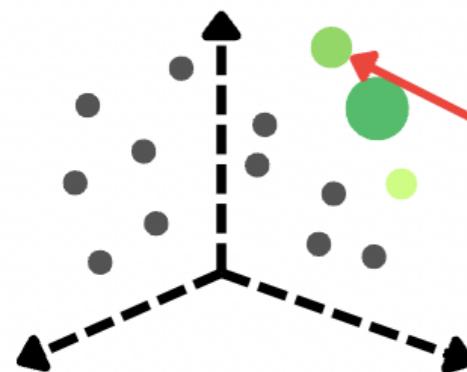
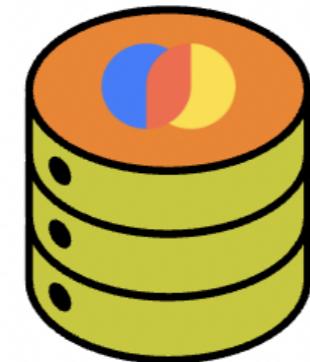


Embedding Model

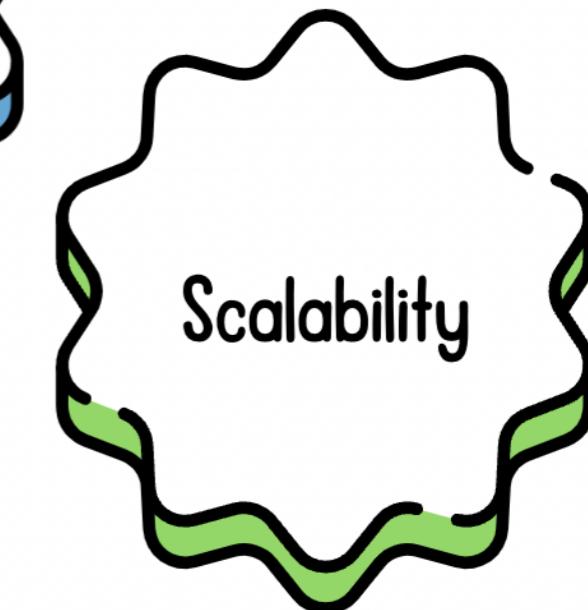
[-0.829647...]



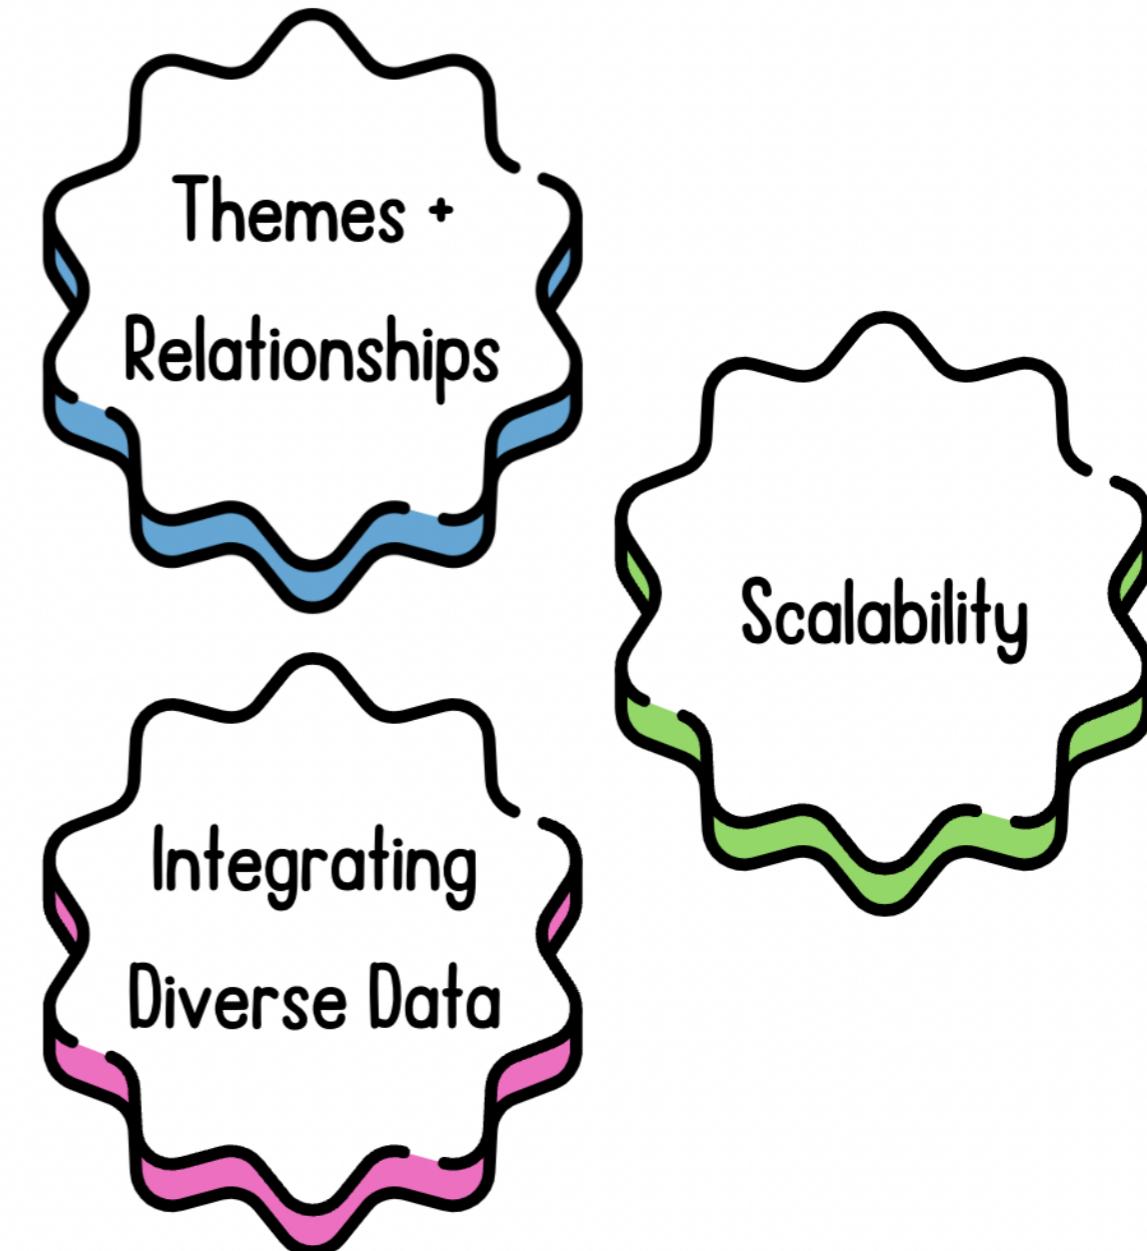
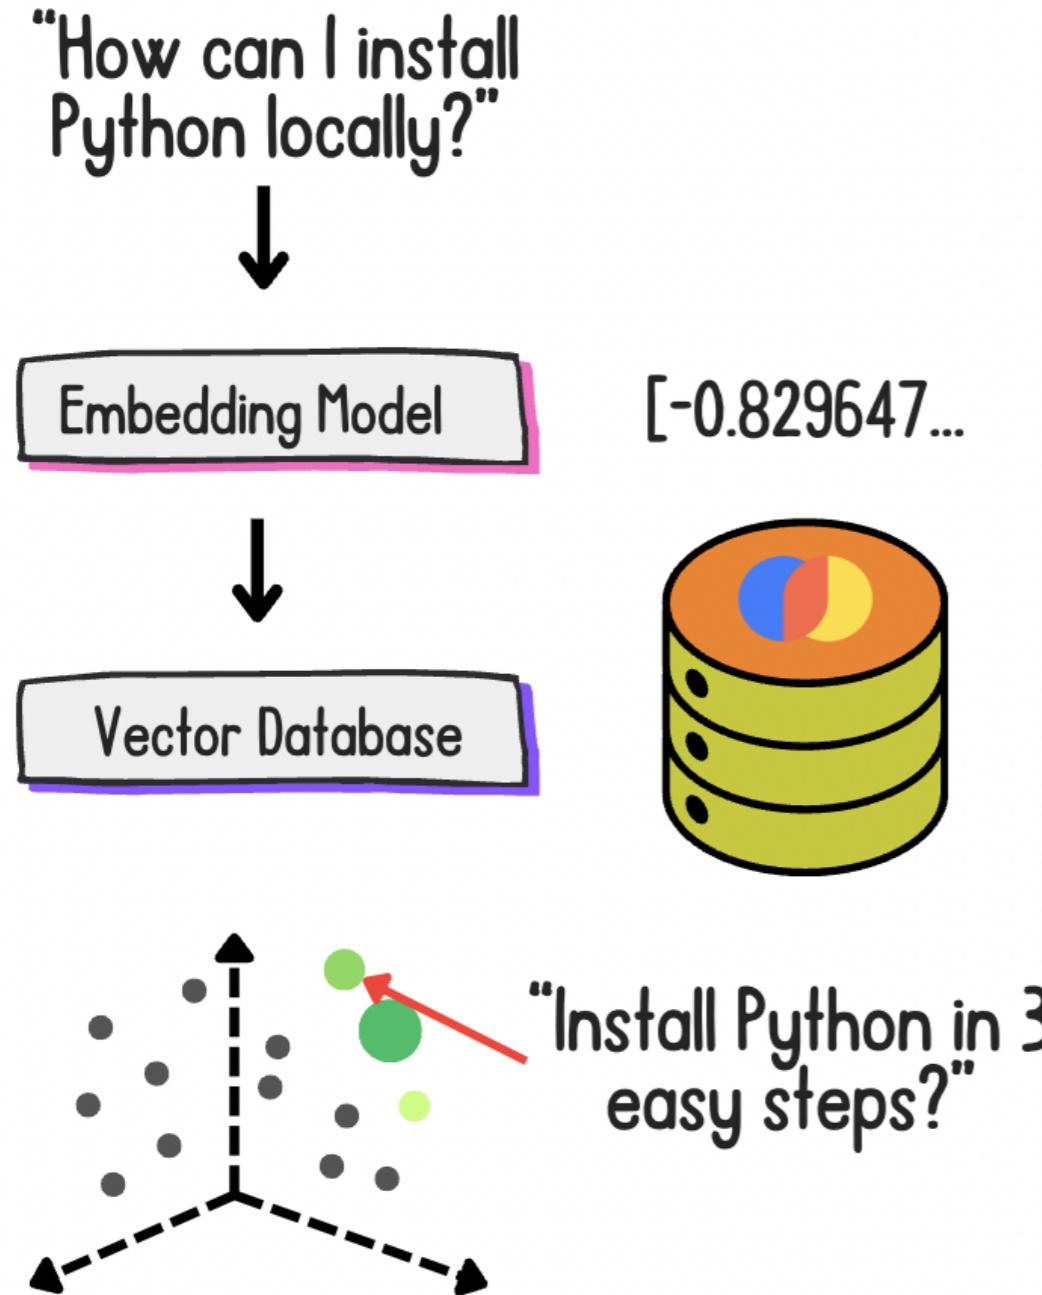
Vector Database



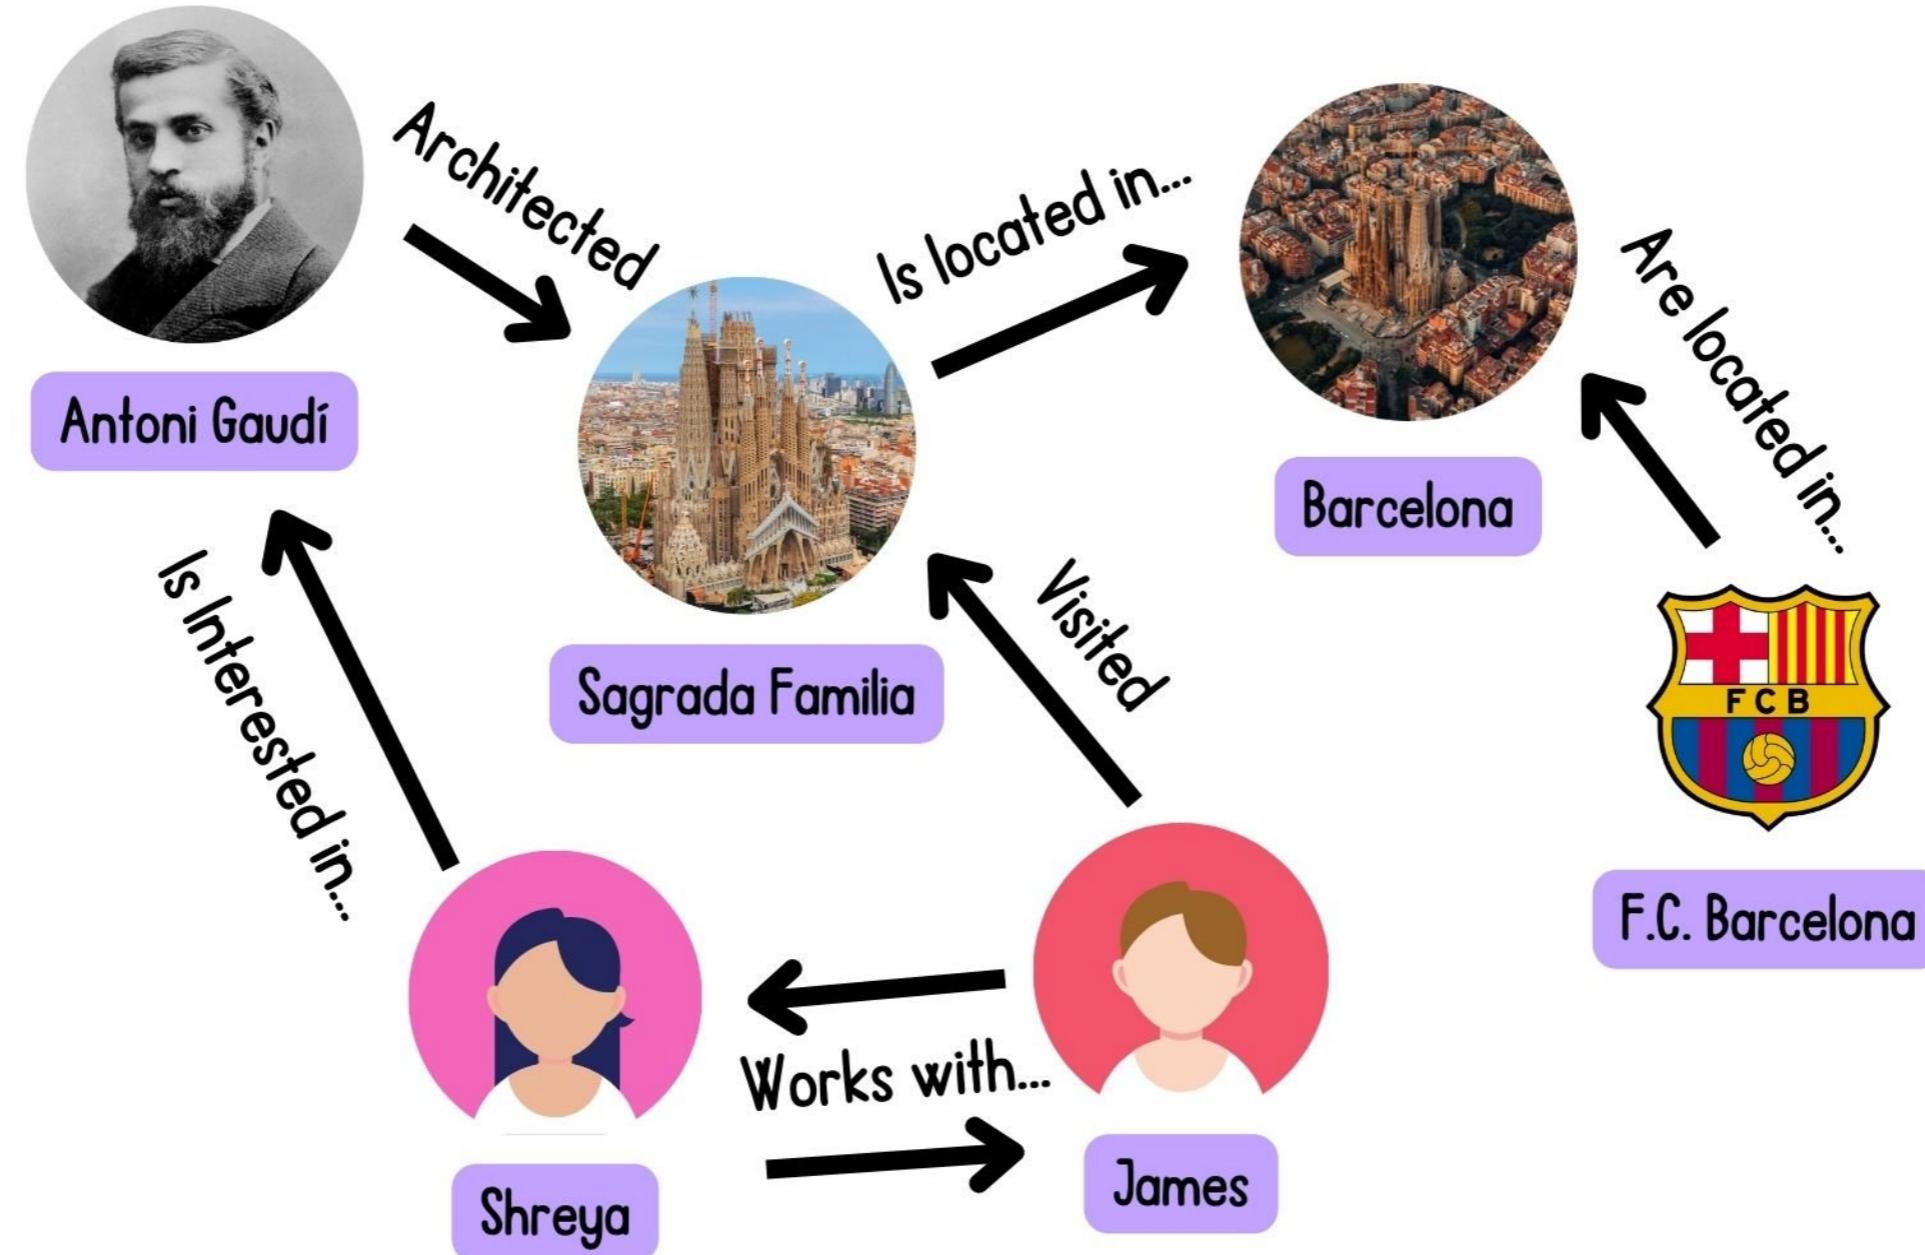
"Install Python in 3  
easy steps?"



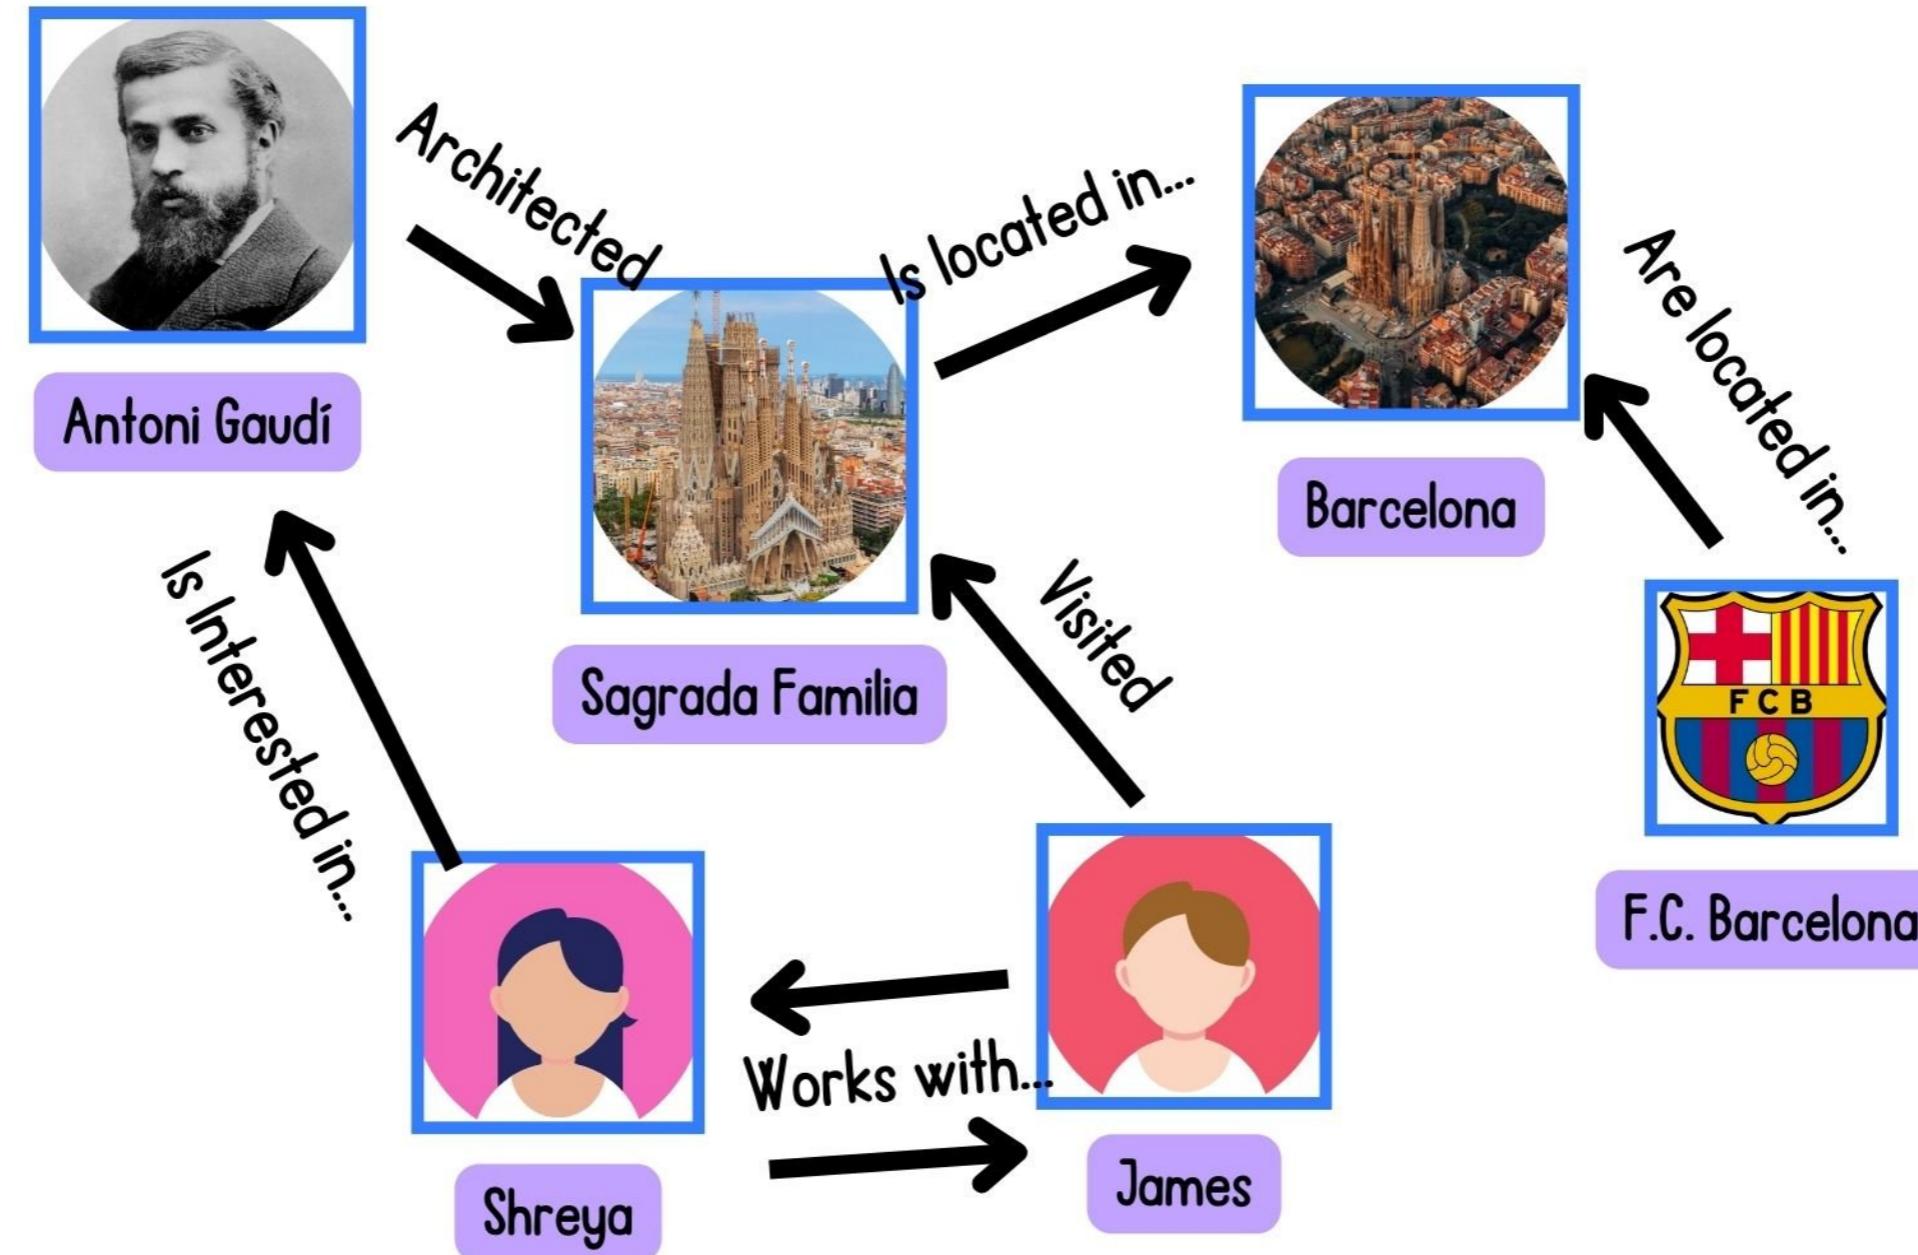
# Vector RAG limitations



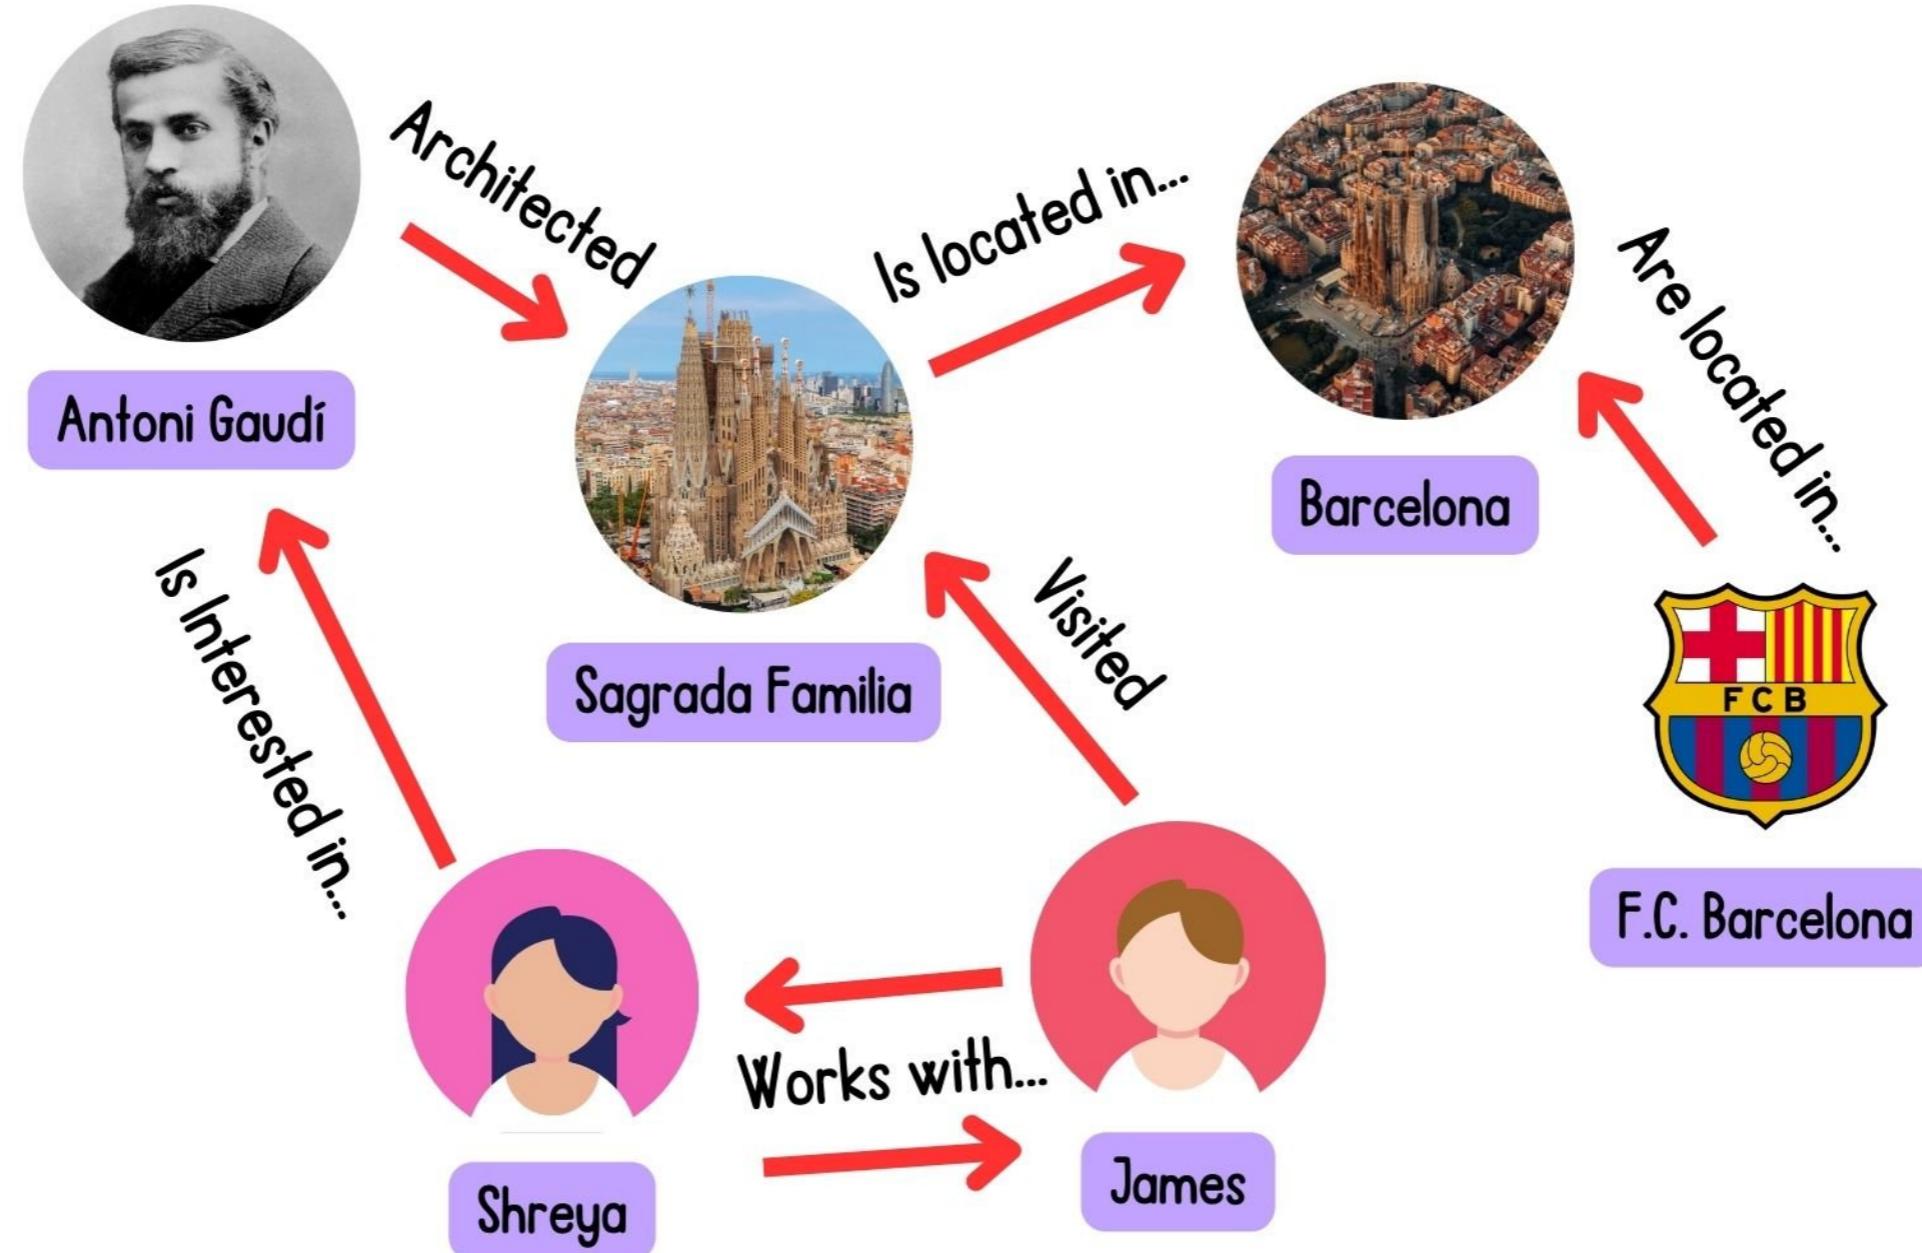
# Graph databases



# Graph databases - nodes



# Graph databases - edges

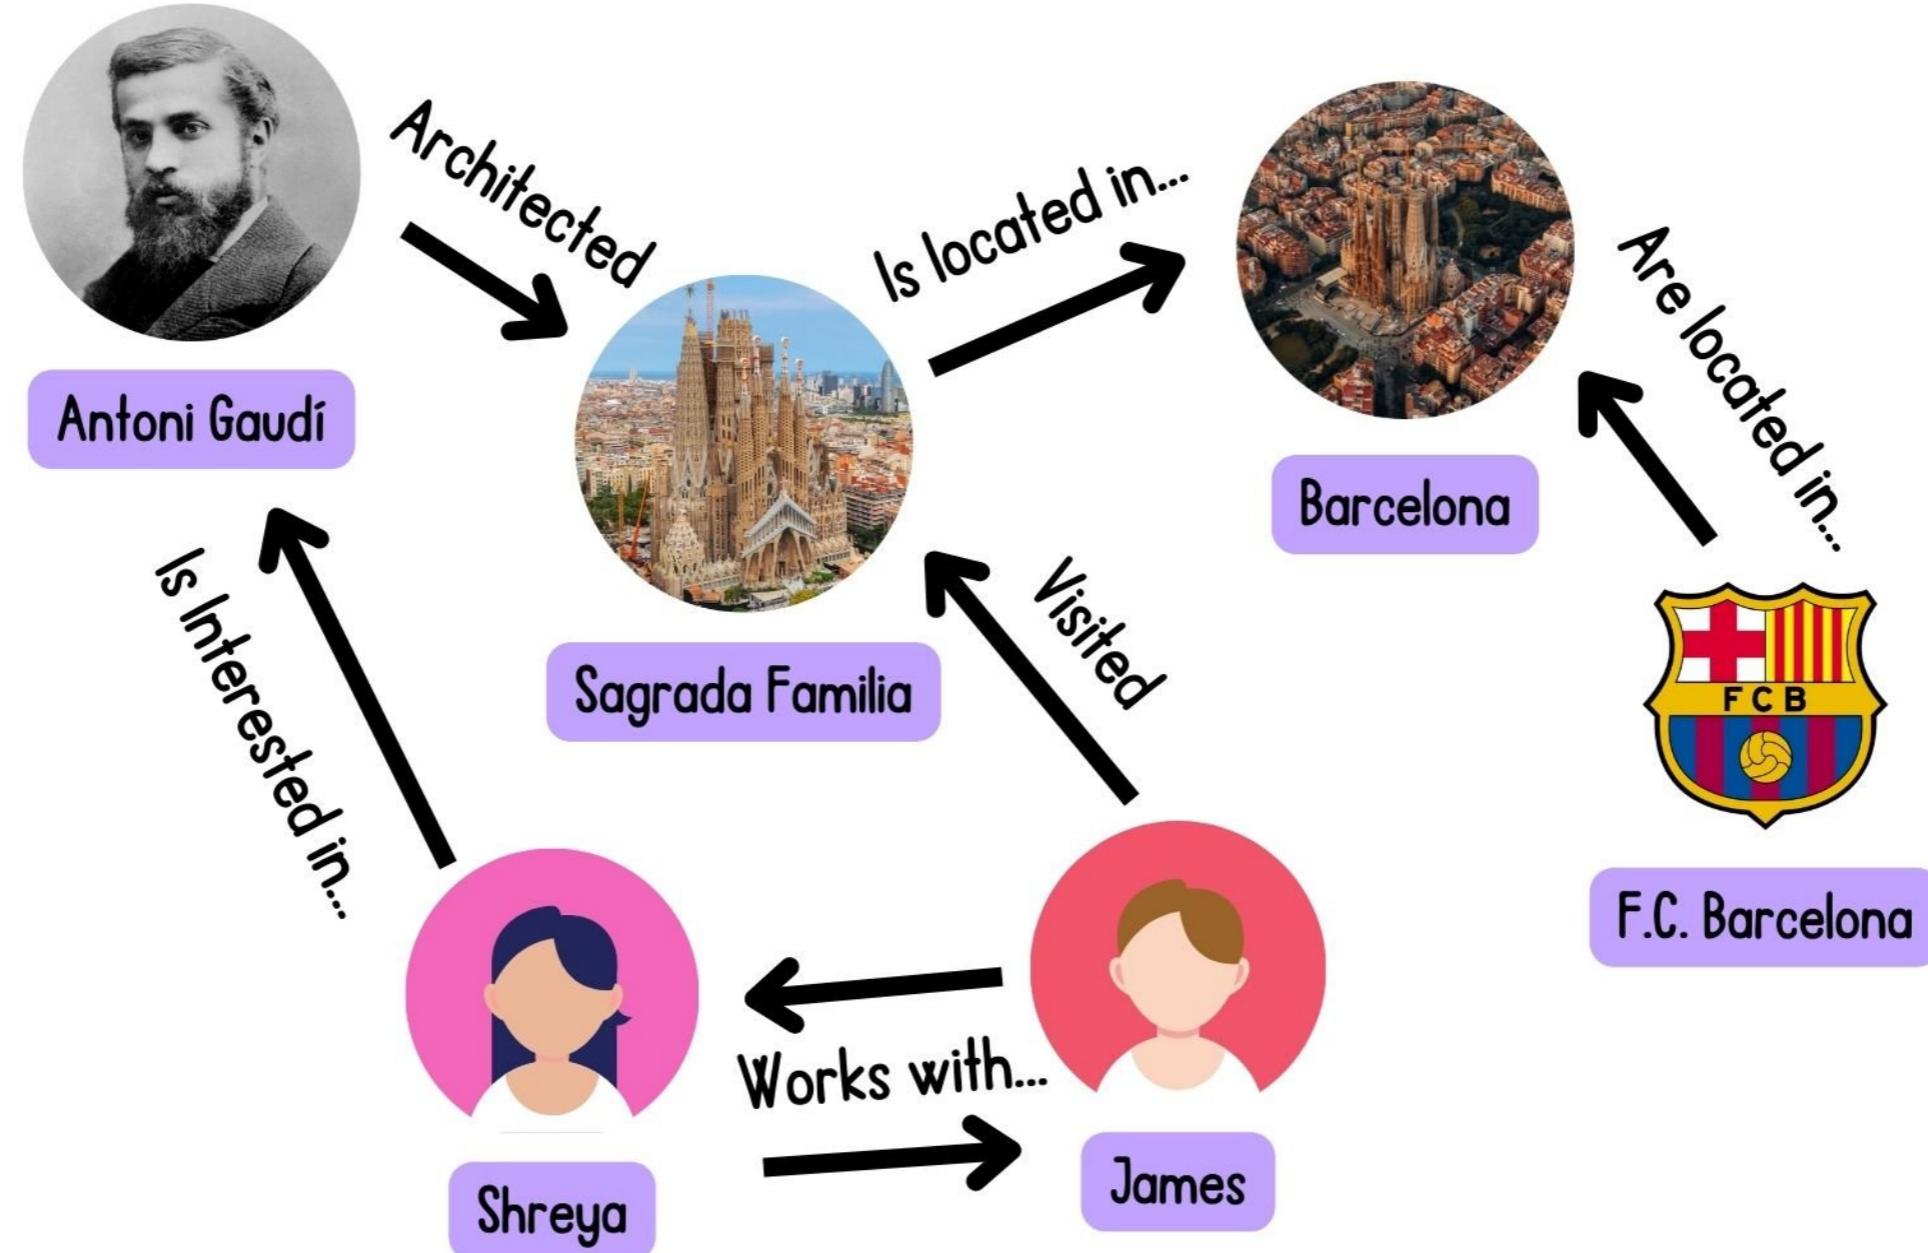


# Neo4j graph databases

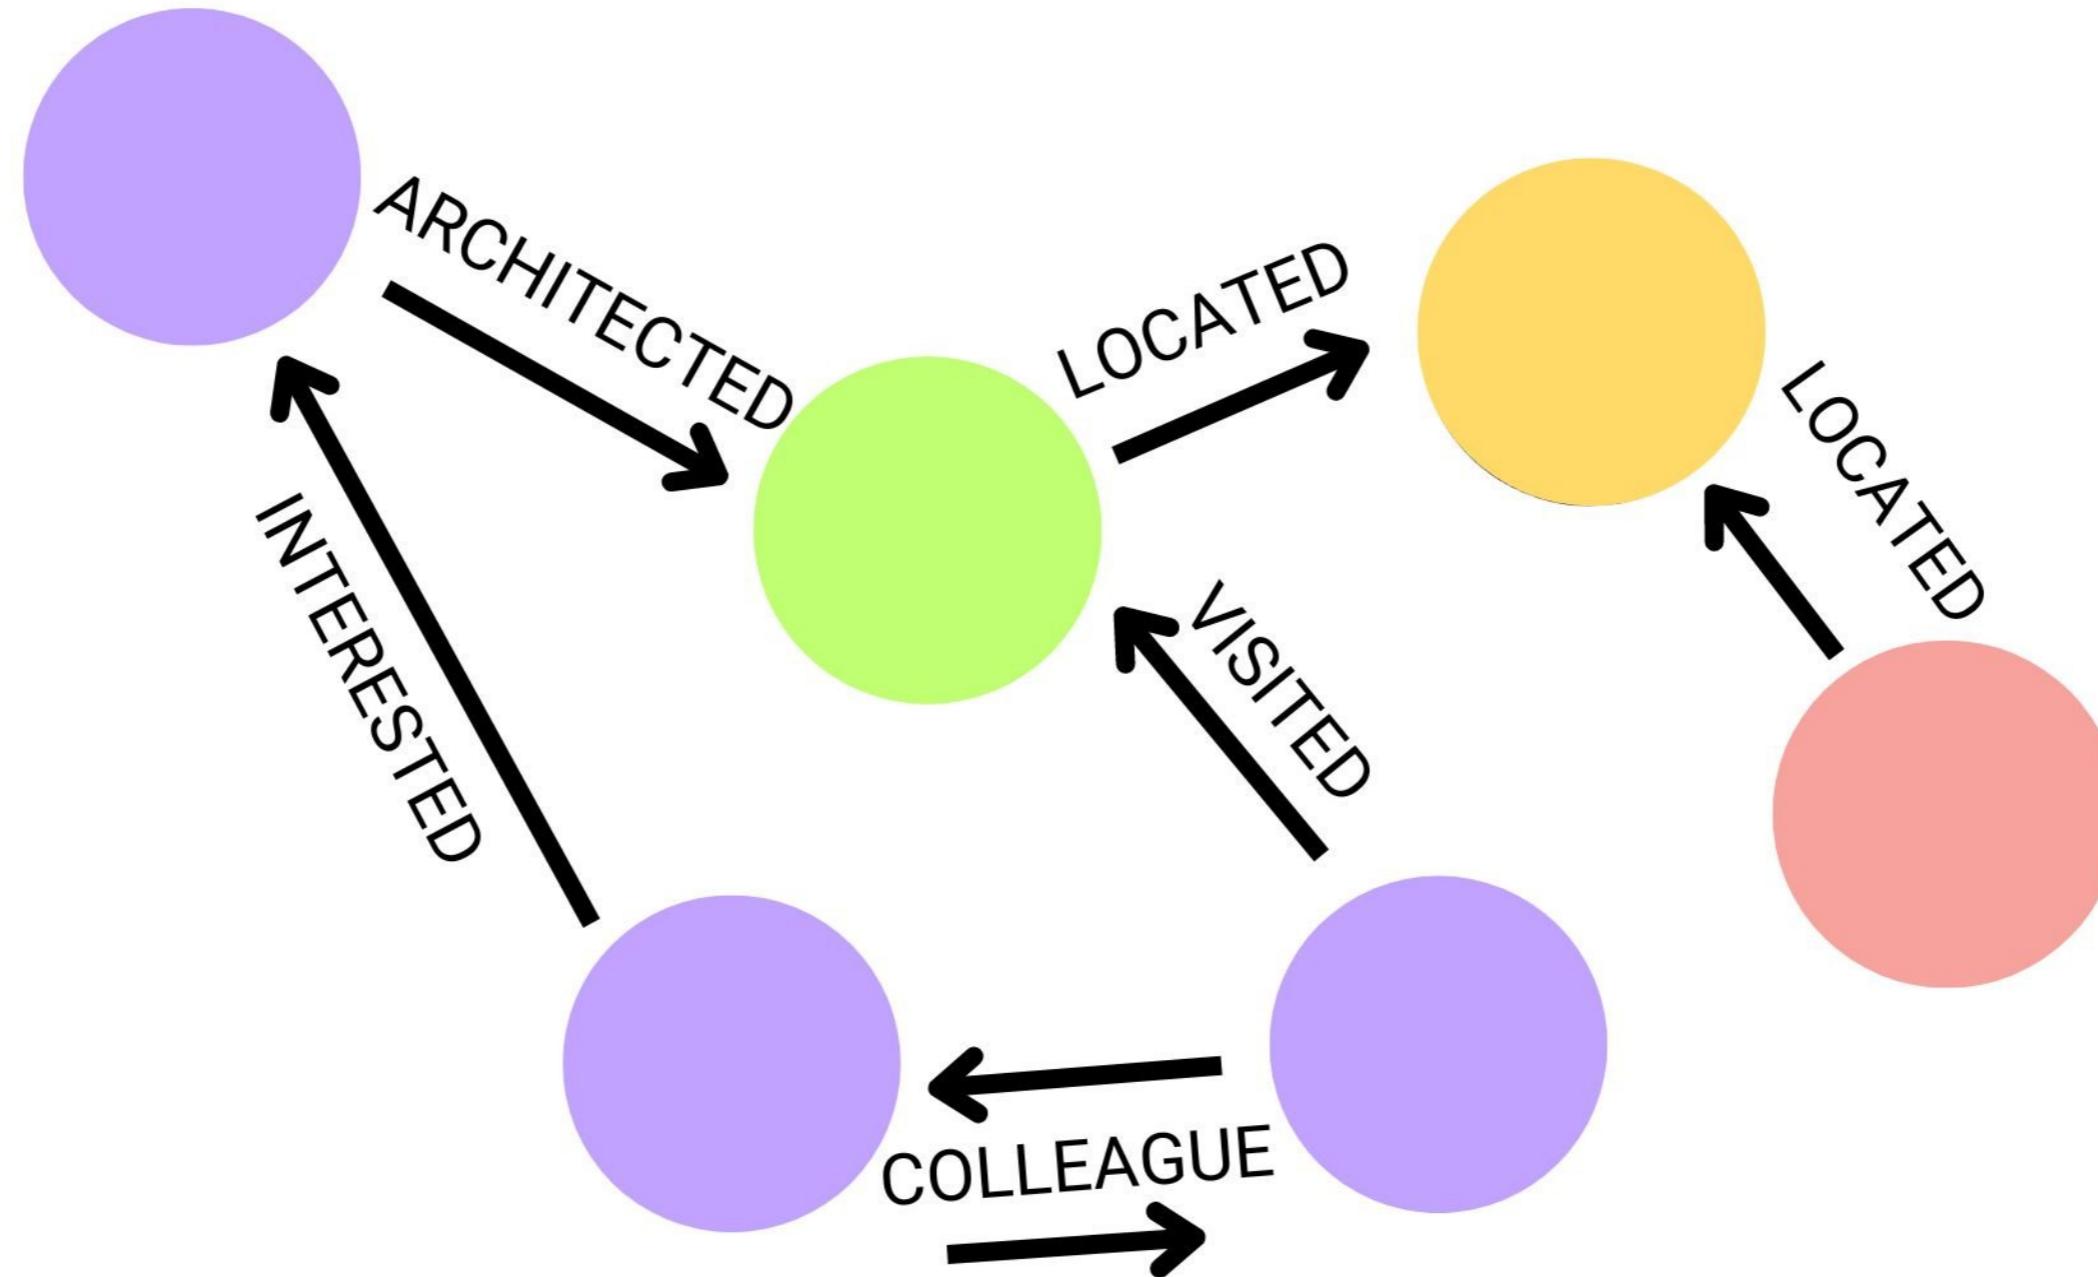


LangChain

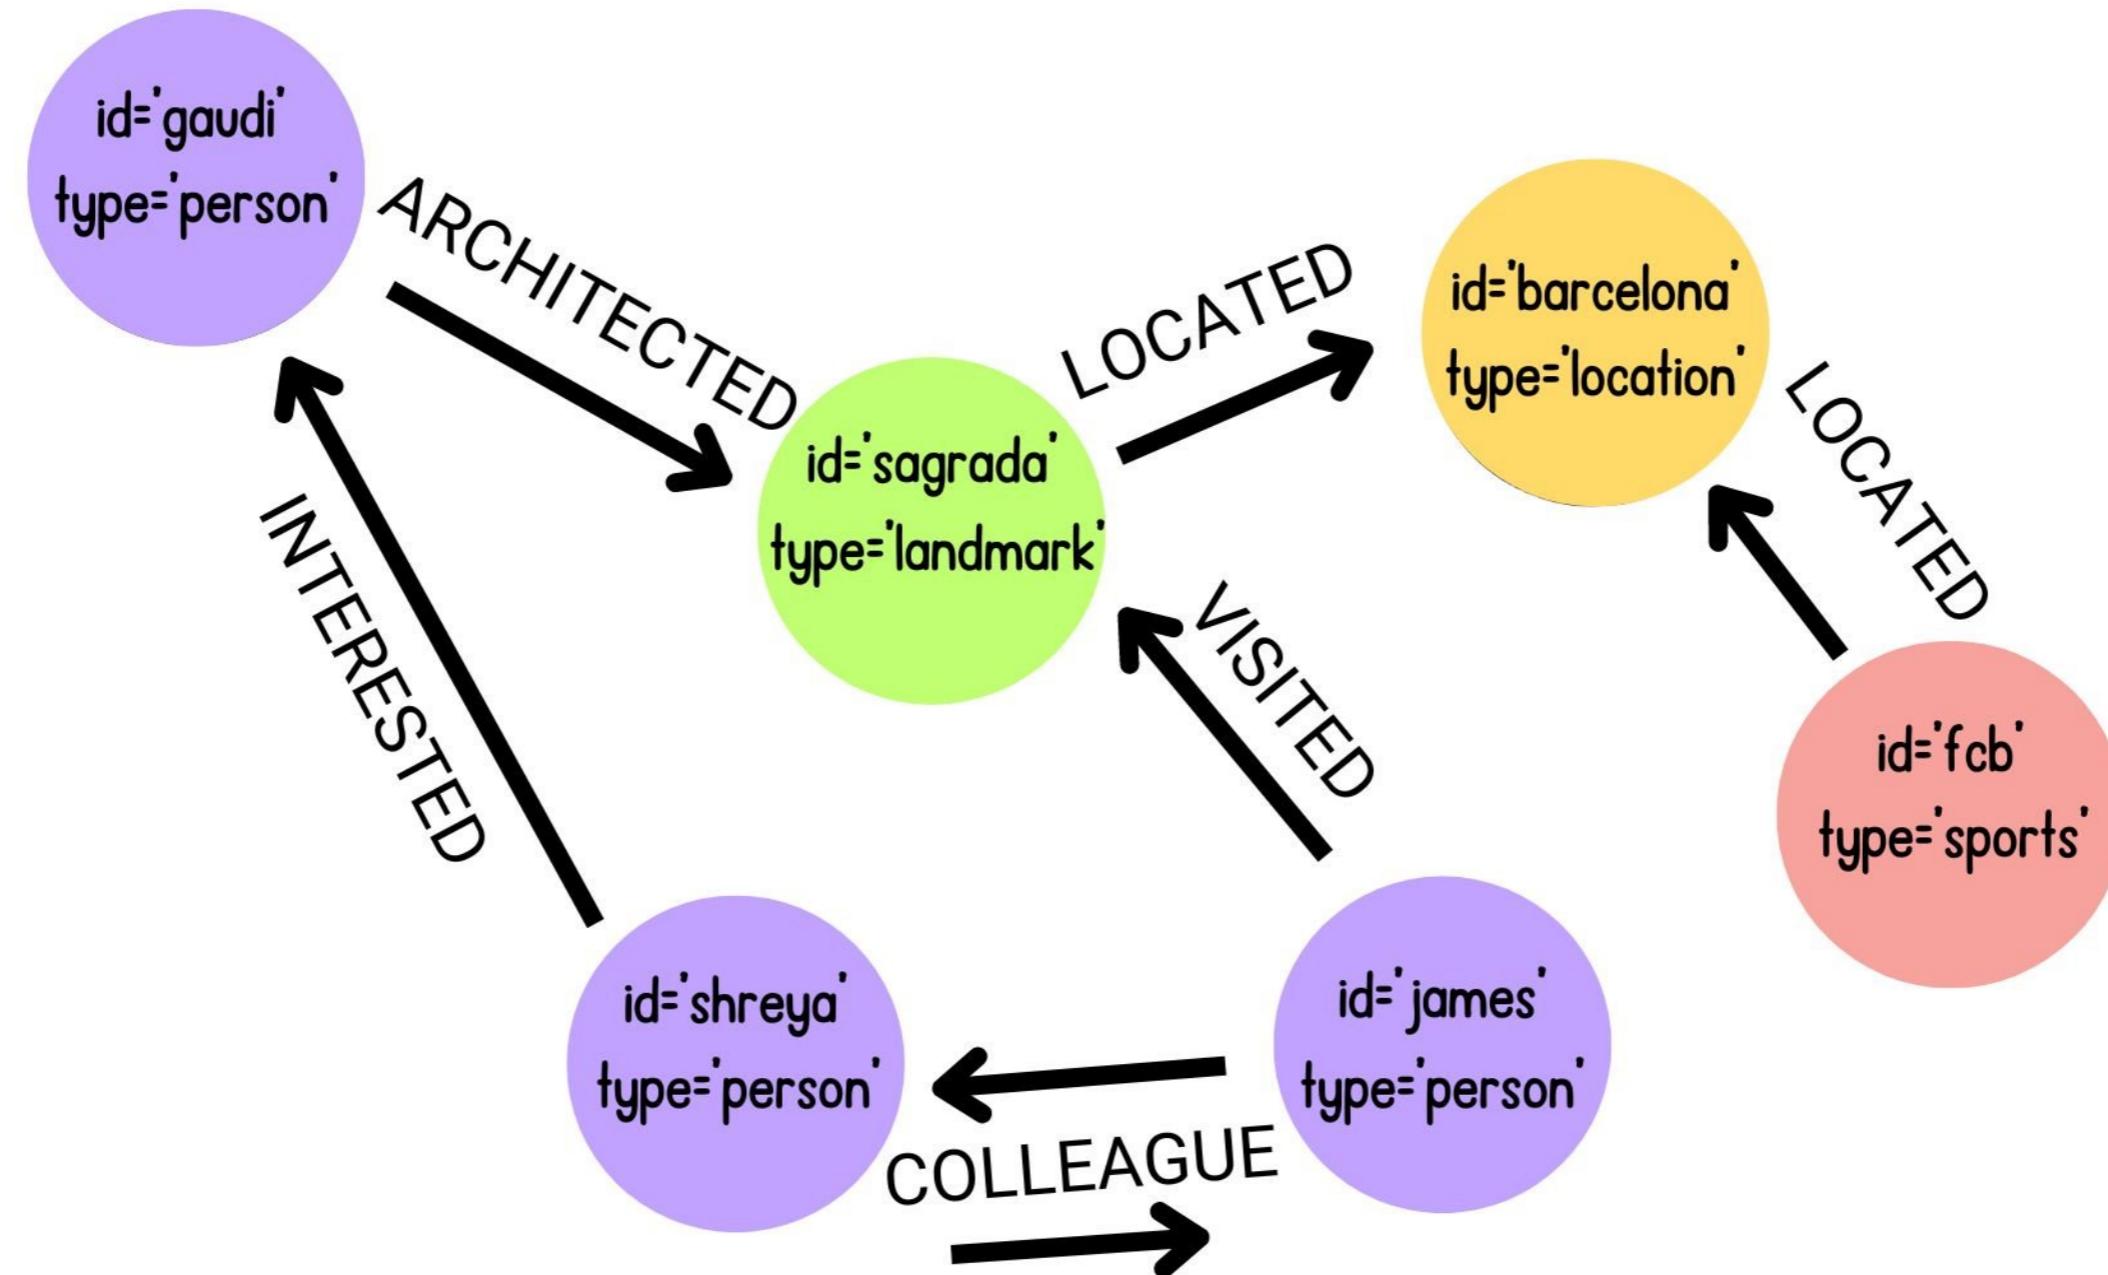
# From graphics to graphs...



# From graphics to graphs...



# From graphics to graphs...



# Loading and chunking Wikipedia pages

```
from langchain_community.document_loaders import WikipediaLoader
from langchain_text_splitters import TokenTextSplitter

raw_documents = WikipediaLoader(query="large language model").load()
text_splitter = TokenTextSplitter(chunk_size=100, chunk_overlap=20)
documents = text_splitter.split_documents(raw_documents[:3])

print(documents[0])
```

```
page_content='A large language model (LLM) is a computational model capable of...'
metadata={'title': 'Large language model',
          'summary': "A large language model (LLM) is...",
          'source': 'https://en.wikipedia.org/wiki/Large_language_model'}
```

# From text to graphs!

```
from langchain_openai import ChatOpenAI
from langchain_experimental.graph_transformers import LLMGraphTransformer

llm = ChatOpenAI(api_key="...", temperature=0, model_name="gpt-4o-mini")
llm_transformer = LLMGraphTransformer(llm=llm)

graph_documents = llm_transformer.convert_to_graph_documents(documents)
print(graph_documents)
```

# From text to graphs!

```
[GraphDocument(  
    nodes=[  
        Node(id='Llm', type='Computational model'),  
        Node(id='Language Generation', type='Concept'),  
        Node(id='Natural Language Processing Tasks', type='Concept'),  
        Node(id='Llama Family', type='Computational model'),  
        Node(id='Ibm', type='Organization'),  
        ..., Node(id='Bert', type='Computational model')],  
    relationships=[  
        Relationship(source=Node(id='Llm', type='Computational model'),  
                    target=Node(id='Language Generation', type='Concept'),  
                    type='CAPABLE_OF'),  
        ...])]
```

# **Let's practice!**

**RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN**

# Storing and querying documents

RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN



**Meri Nova**  
Machine Learning Engineer

# Instantiating the Neo4j database

```
from langchain_community.graphs import Neo4jGraph
```

```
graph = Neo4jGraph(url="bolt://localhost:7687", username="neo4j", password="...")
```

```
import os
```

```
url = os.environ["NEO4J_URI"]
```

```
user = os.environ["NEO4J_USERNAME"]
```

```
password = os.environ["NEO4J_PASSWORD"]
```

```
graph = Neo4jGraph(url=url, username=user, password=password)
```

<sup>1</sup> <https://neo4j.com/download/>

# Storing graph documents

```
from langchain_experimental.graph_transformers import LLMGraphTransformer

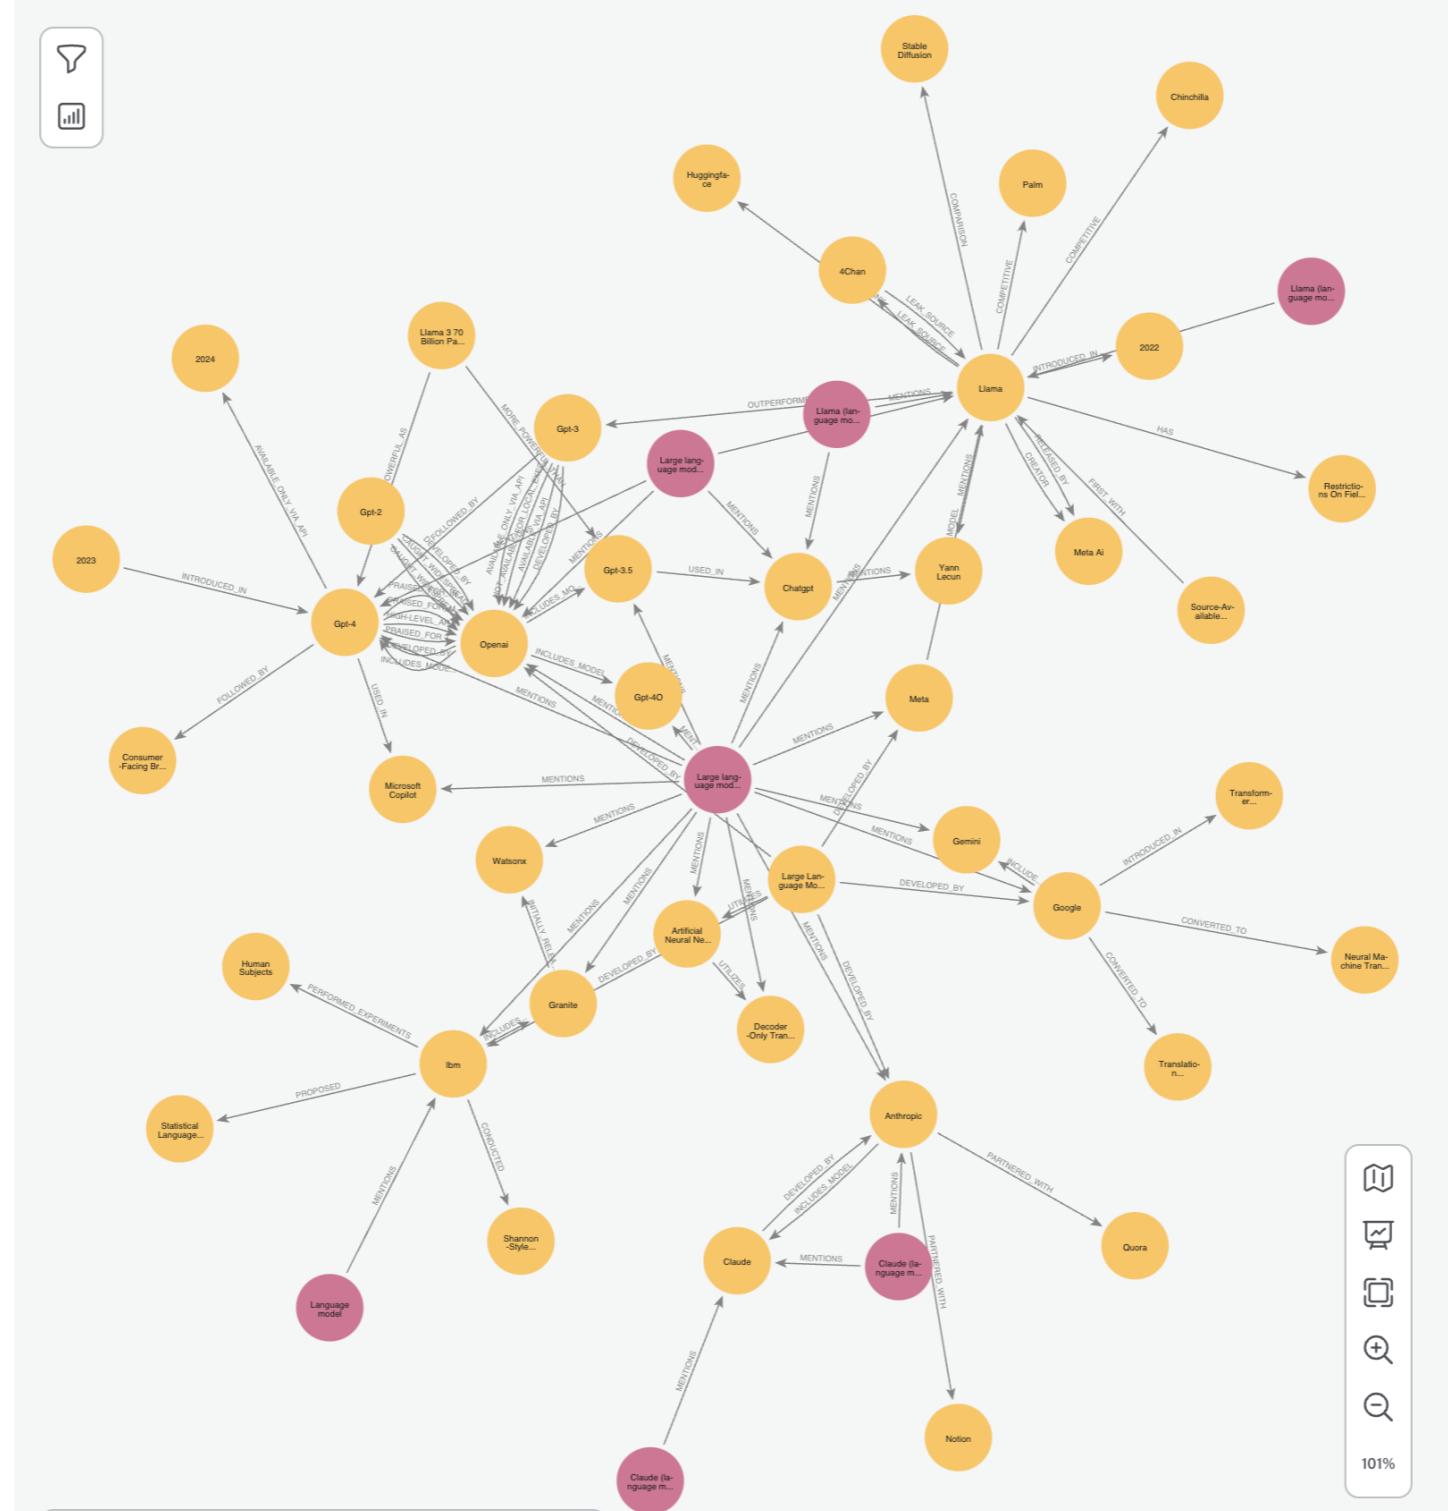
llm = ChatOpenAI(api_key="...", temperature=0, model="gpt-4o-mini")
llm_transformer = LLMGraphTransformer(llm=llm)

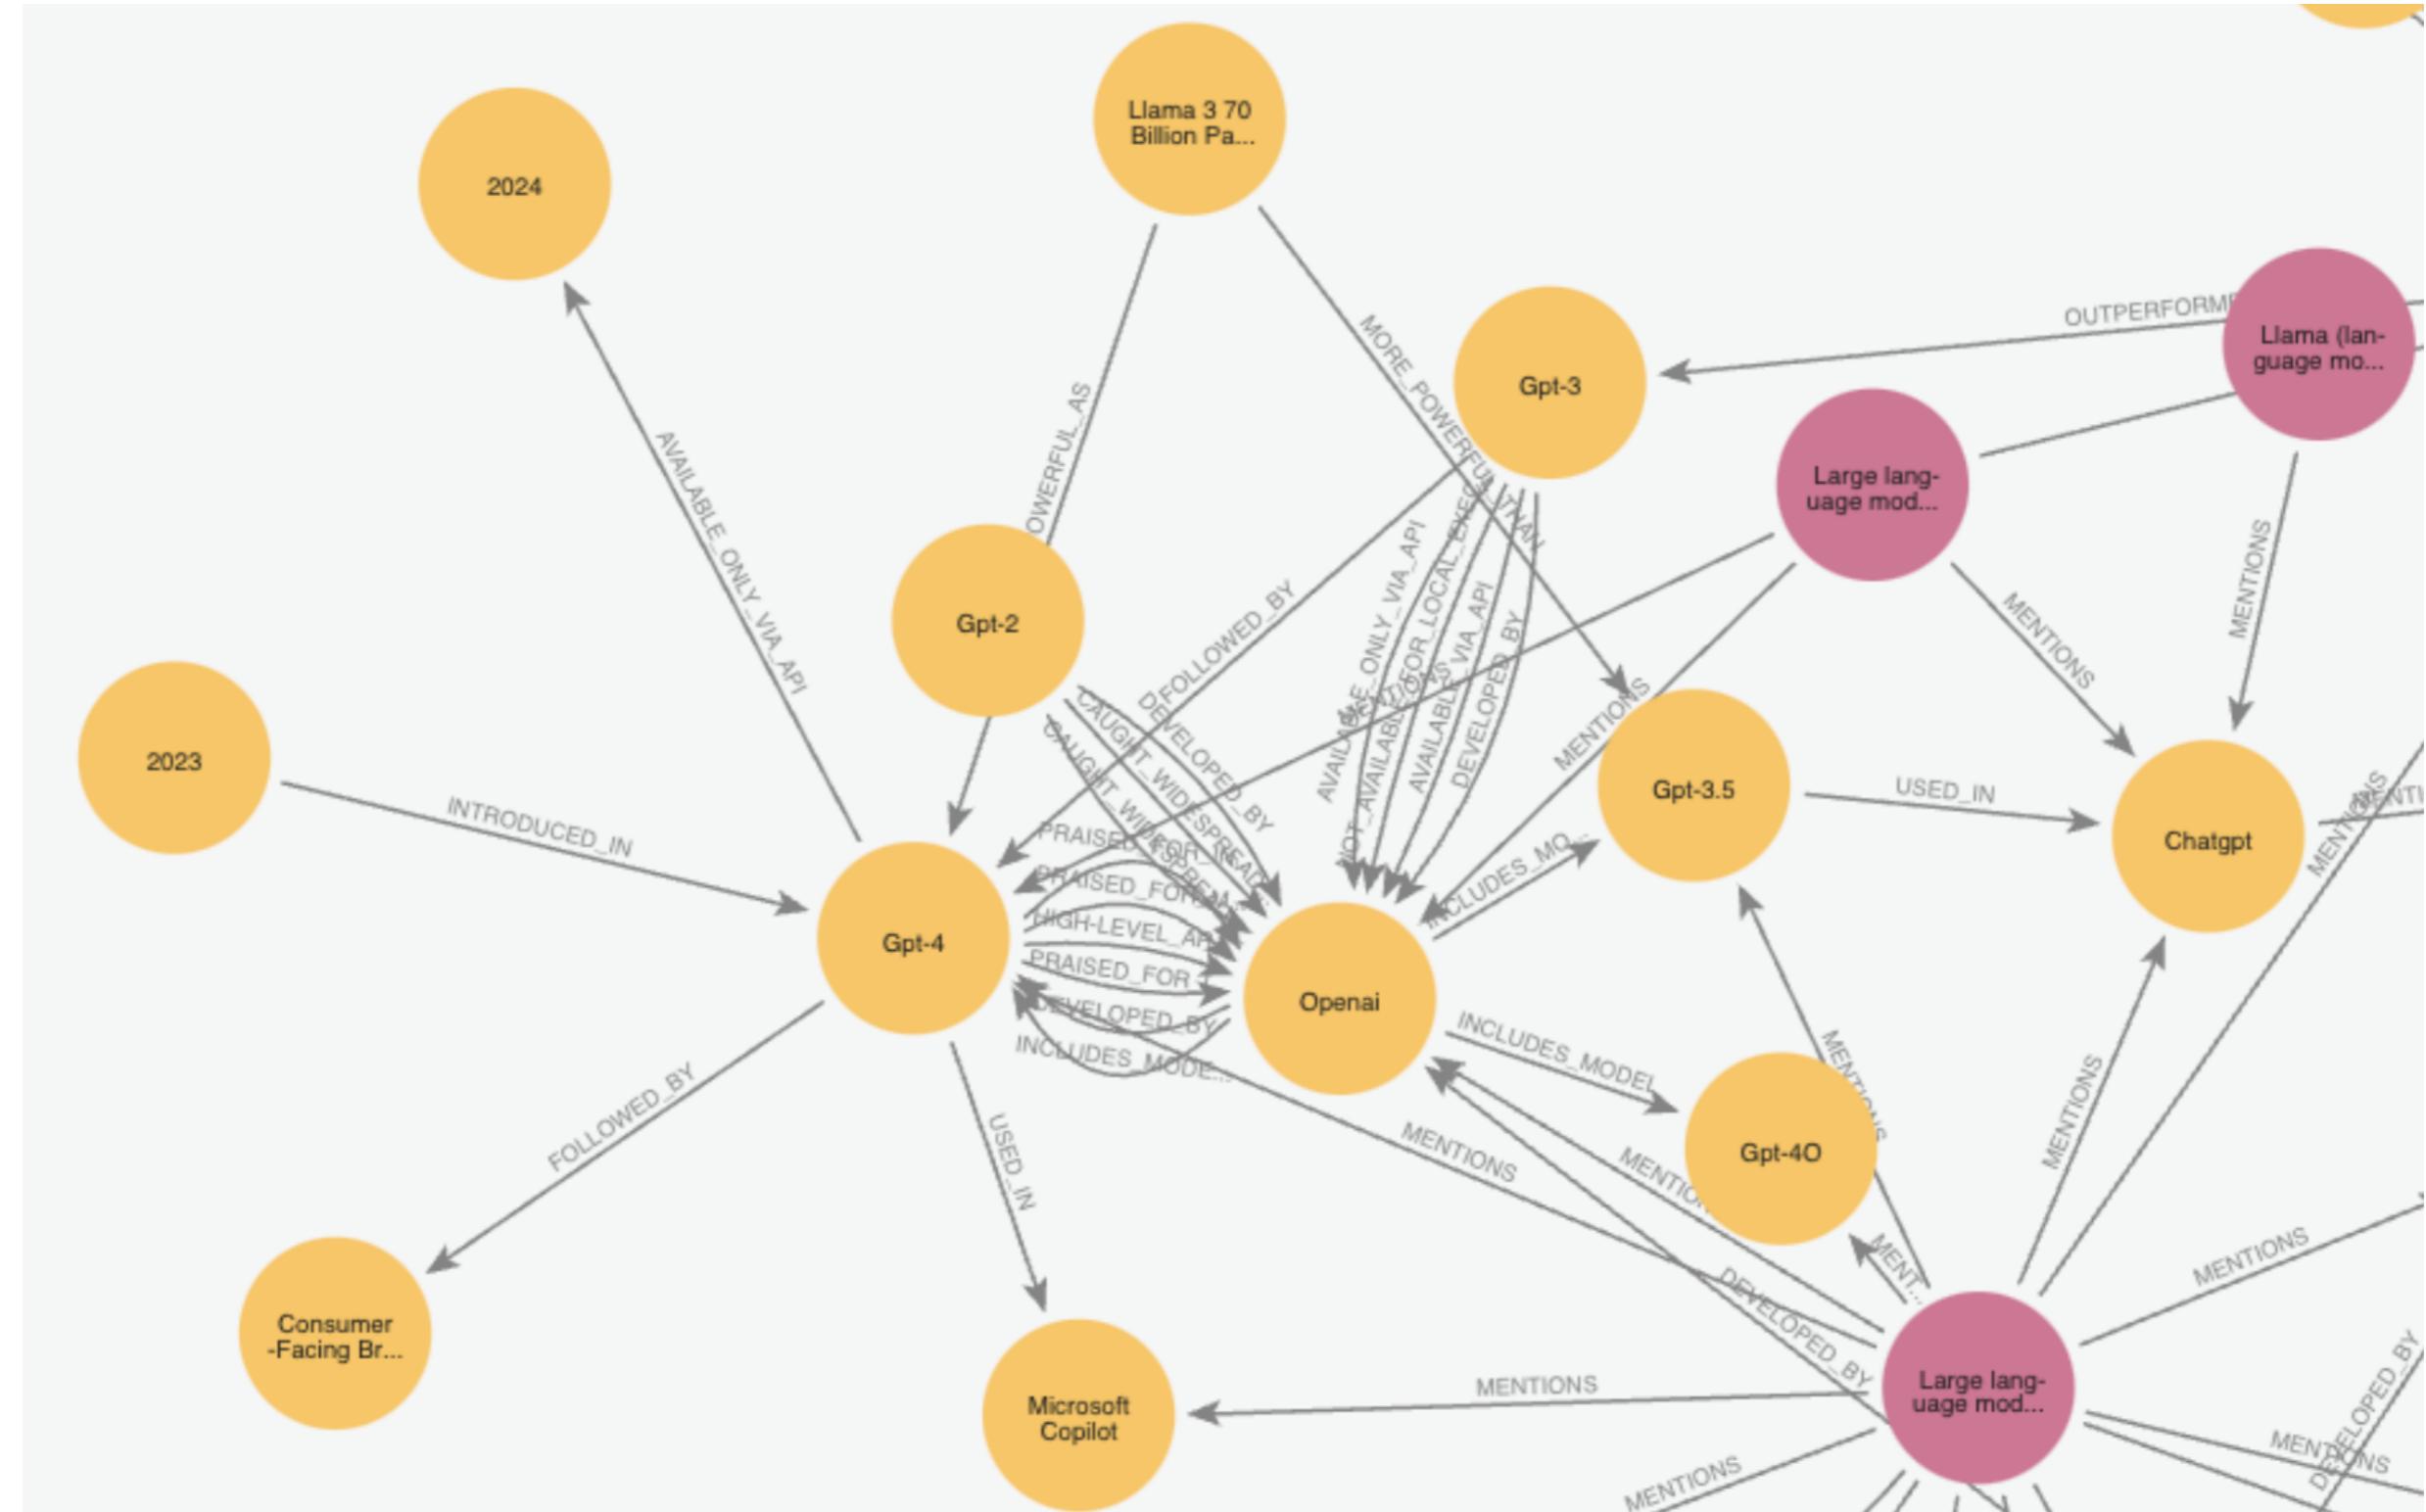
graph_documents = llm_transformer.convert_to_graph_documents(documents)
```

# Storing graph documents

```
graph.add_graph_documents(  
    graph_documents,  
    include_source=True,  
    baseEntityLabel=True  
)
```

- `include_source=True` : link nodes to source documents with `MENTIONS` edge
- `baseEntityLabel=True` : add `__Entity__` label to each node





# Database schema

```
print(graph.get_schema)
```

Node properties:

Concept {id: STRING}

Architecture {id: STRING}

Organization {id: STRING}

Event {id: STRING}

Paper {id: STRING}

The relationships:

(:Concept)-[:DEVELOPED\_BY]->(:Person)

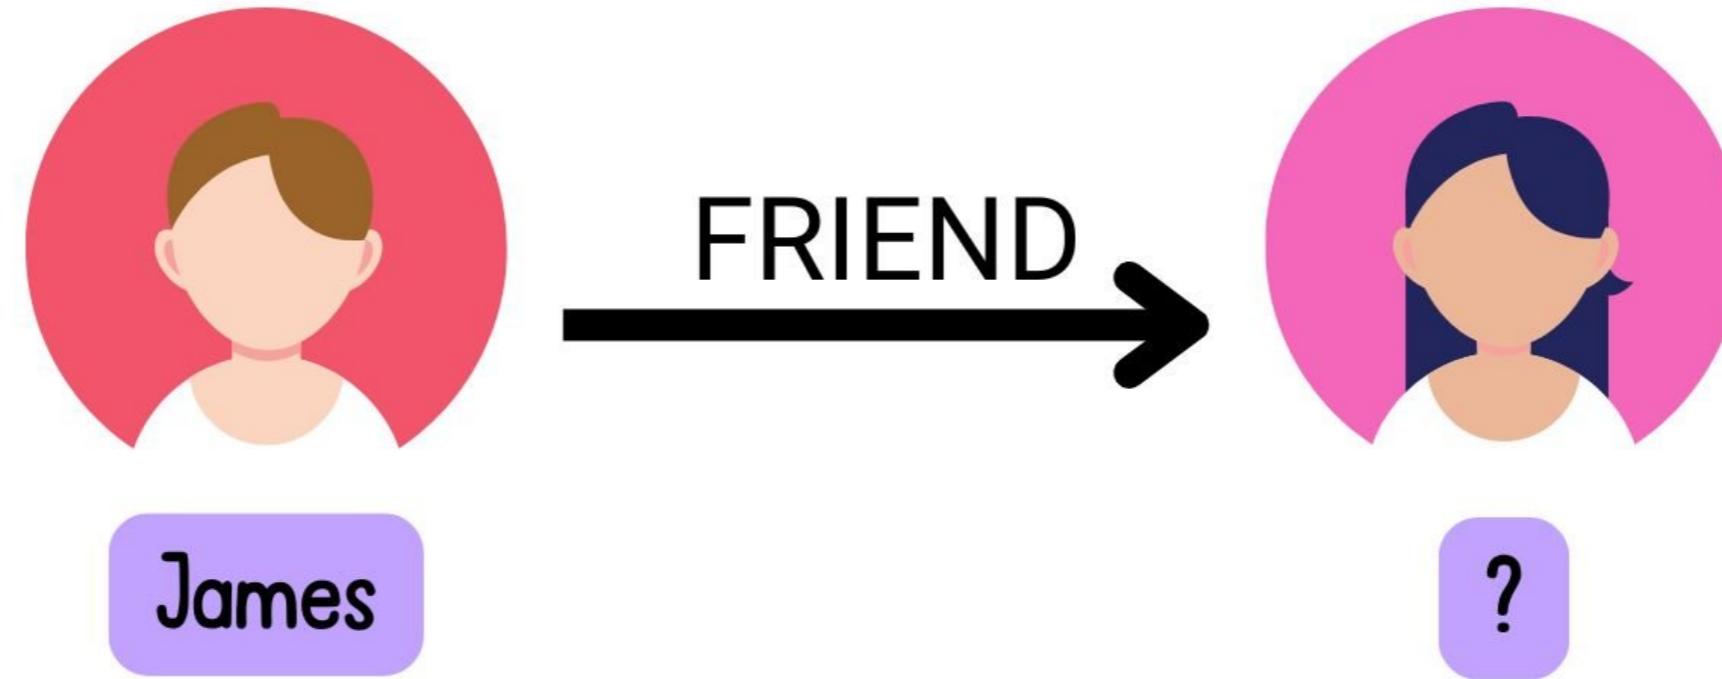
(:Architecture)-[:BASED\_ON]->(:Concept)

(:Organization)-[:PROPOSED]->(:Concept)

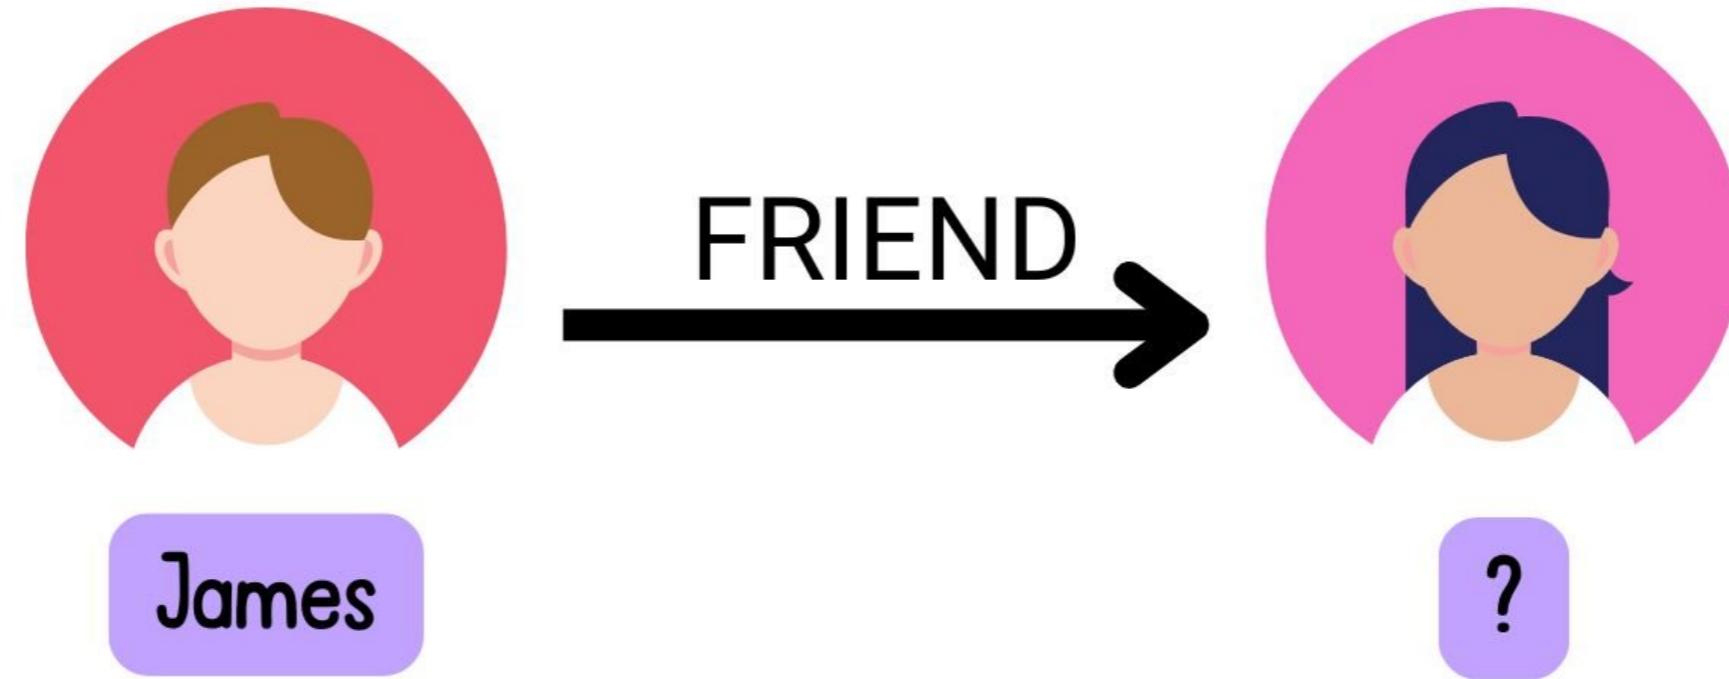
(:Document)-[:MENTIONS]->(:Event)

(:Paper)-[:BASED\_ON]->(:Concept)

# Querying Neo4j - Cypher Query Language

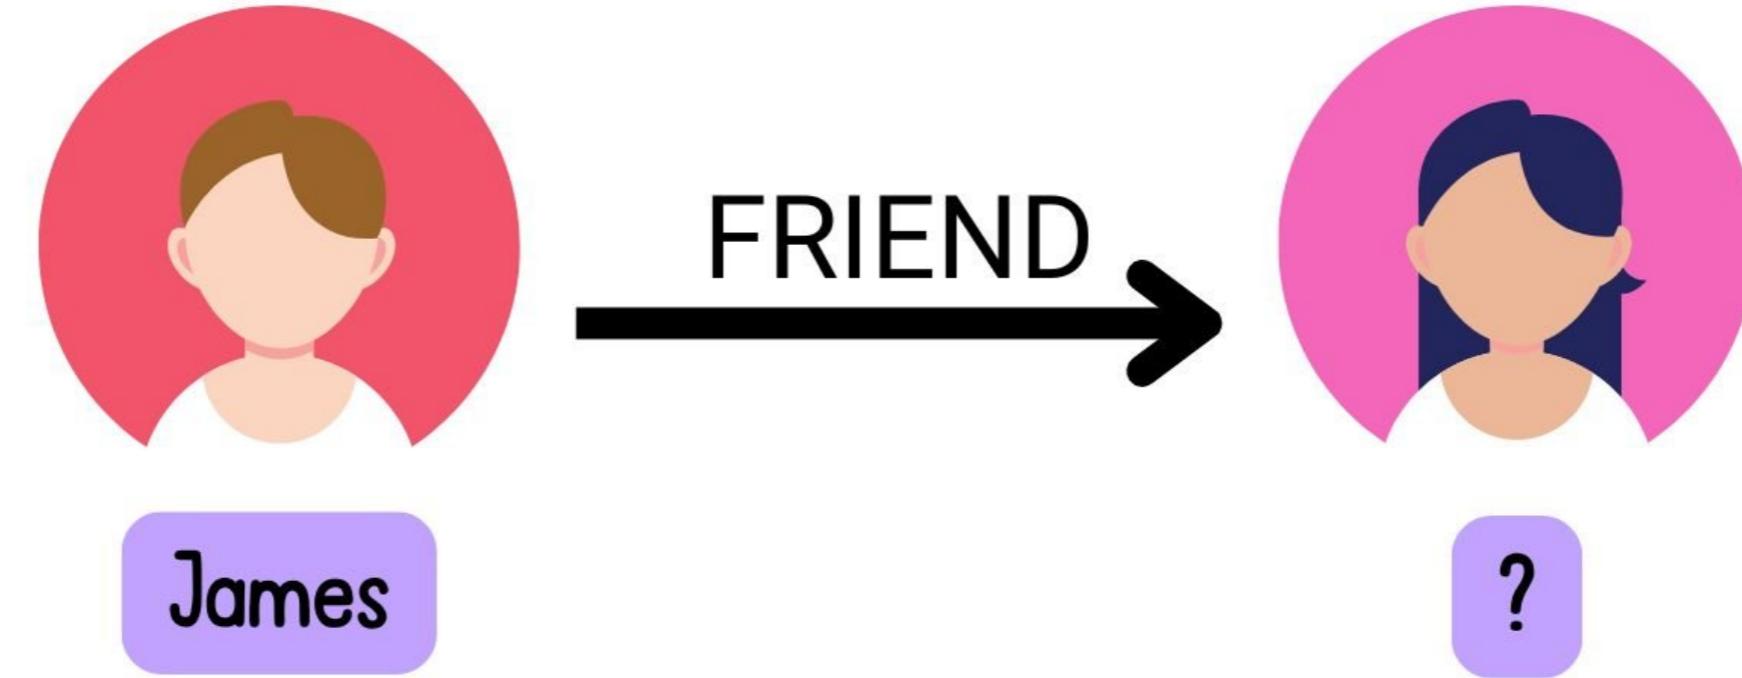


# Querying Neo4j - Cypher Query Language



```
MATCH (james:Person {name: "James"})-[:FRIEND]->(friend) RETURN friend
```

# Querying Neo4j - Cypher Query Language



```
MATCH (james:Person {name: "James"})-[:FRIEND]->(friend) RETURN friend
```

# Querying Neo4j - Cypher Query Language



MATCH (james:Person {name: "James"})-[:FRIEND]->(friend) RETURN friend

Node                      Relationship              Node  
\_\_\_\_\_                    \_\_\_\_\_                    \_\_\_\_\_  
LABEL                      PROPERTY                    VARIABLE

# Querying Neo4j - Cypher Query Language



MATCH (james:Person {name: "James"})-[:FRIEND]->(friend) RETURN friend

Node                      Relationship              Node  
\_\_\_\_\_                    \_\_\_\_\_                    \_\_\_\_\_  
LABEL                      PROPERTY                VARIABLE

# Querying the LLM graph

```
results = graph.query("""  
MATCH (gpt4:Model {id: "Gpt-4"})-[:DEVELOPED_BY]->(org:Organization)  
RETURN org  
""")  
  
print(results)
```

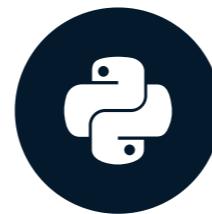
```
[{'org': {'id': 'Openai'}}]
```

# **Let's practice!**

**RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN**

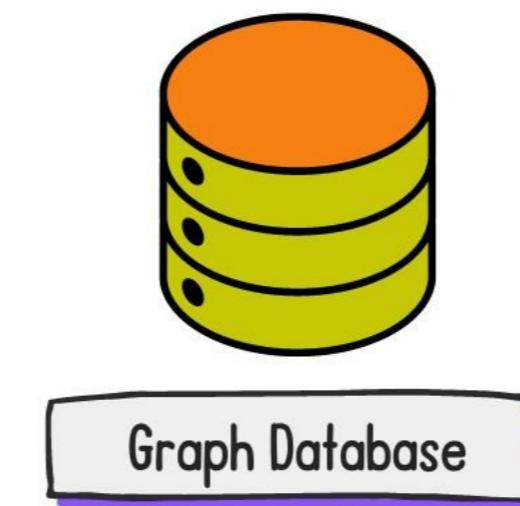
# Creating the Graph RAG chain

RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN



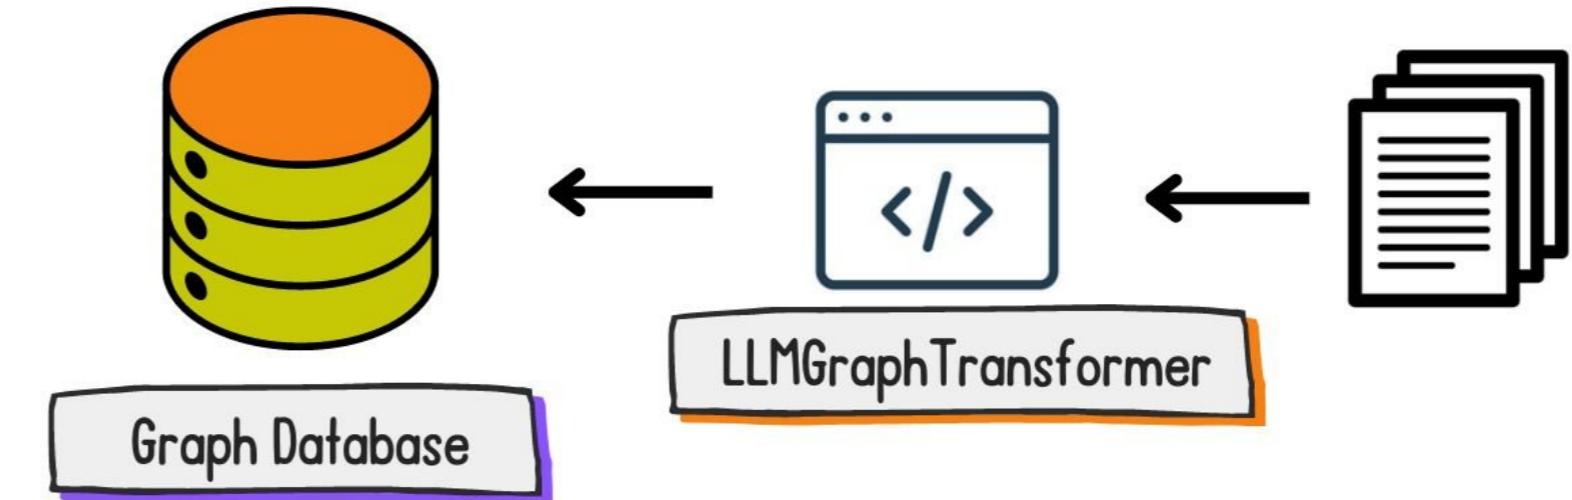
**Meri Nova**  
Machine Learning Engineer

# Building the Graph RAG architecture

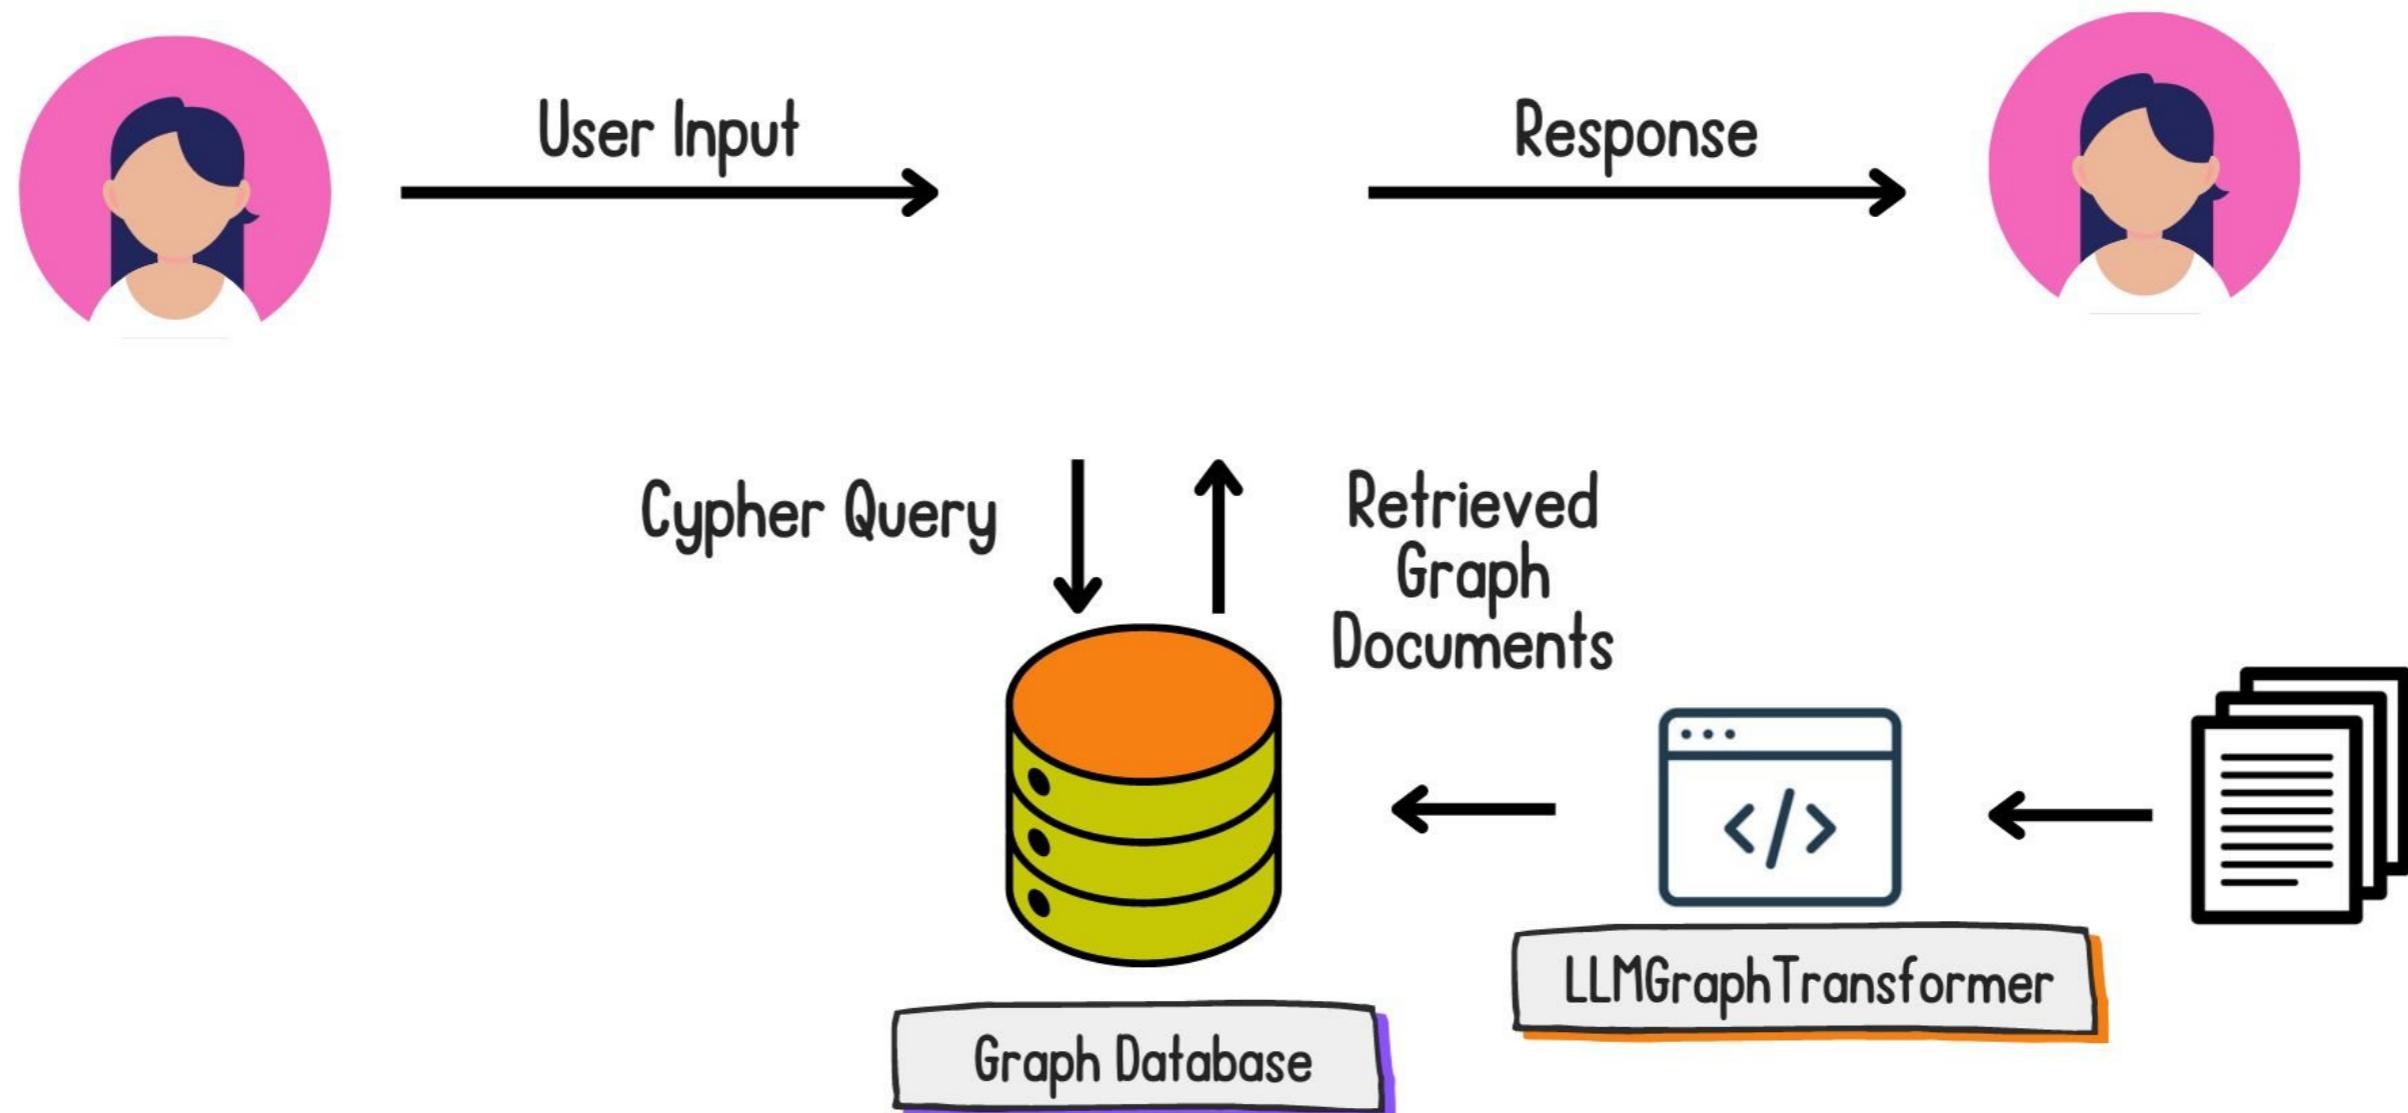


Graph Database

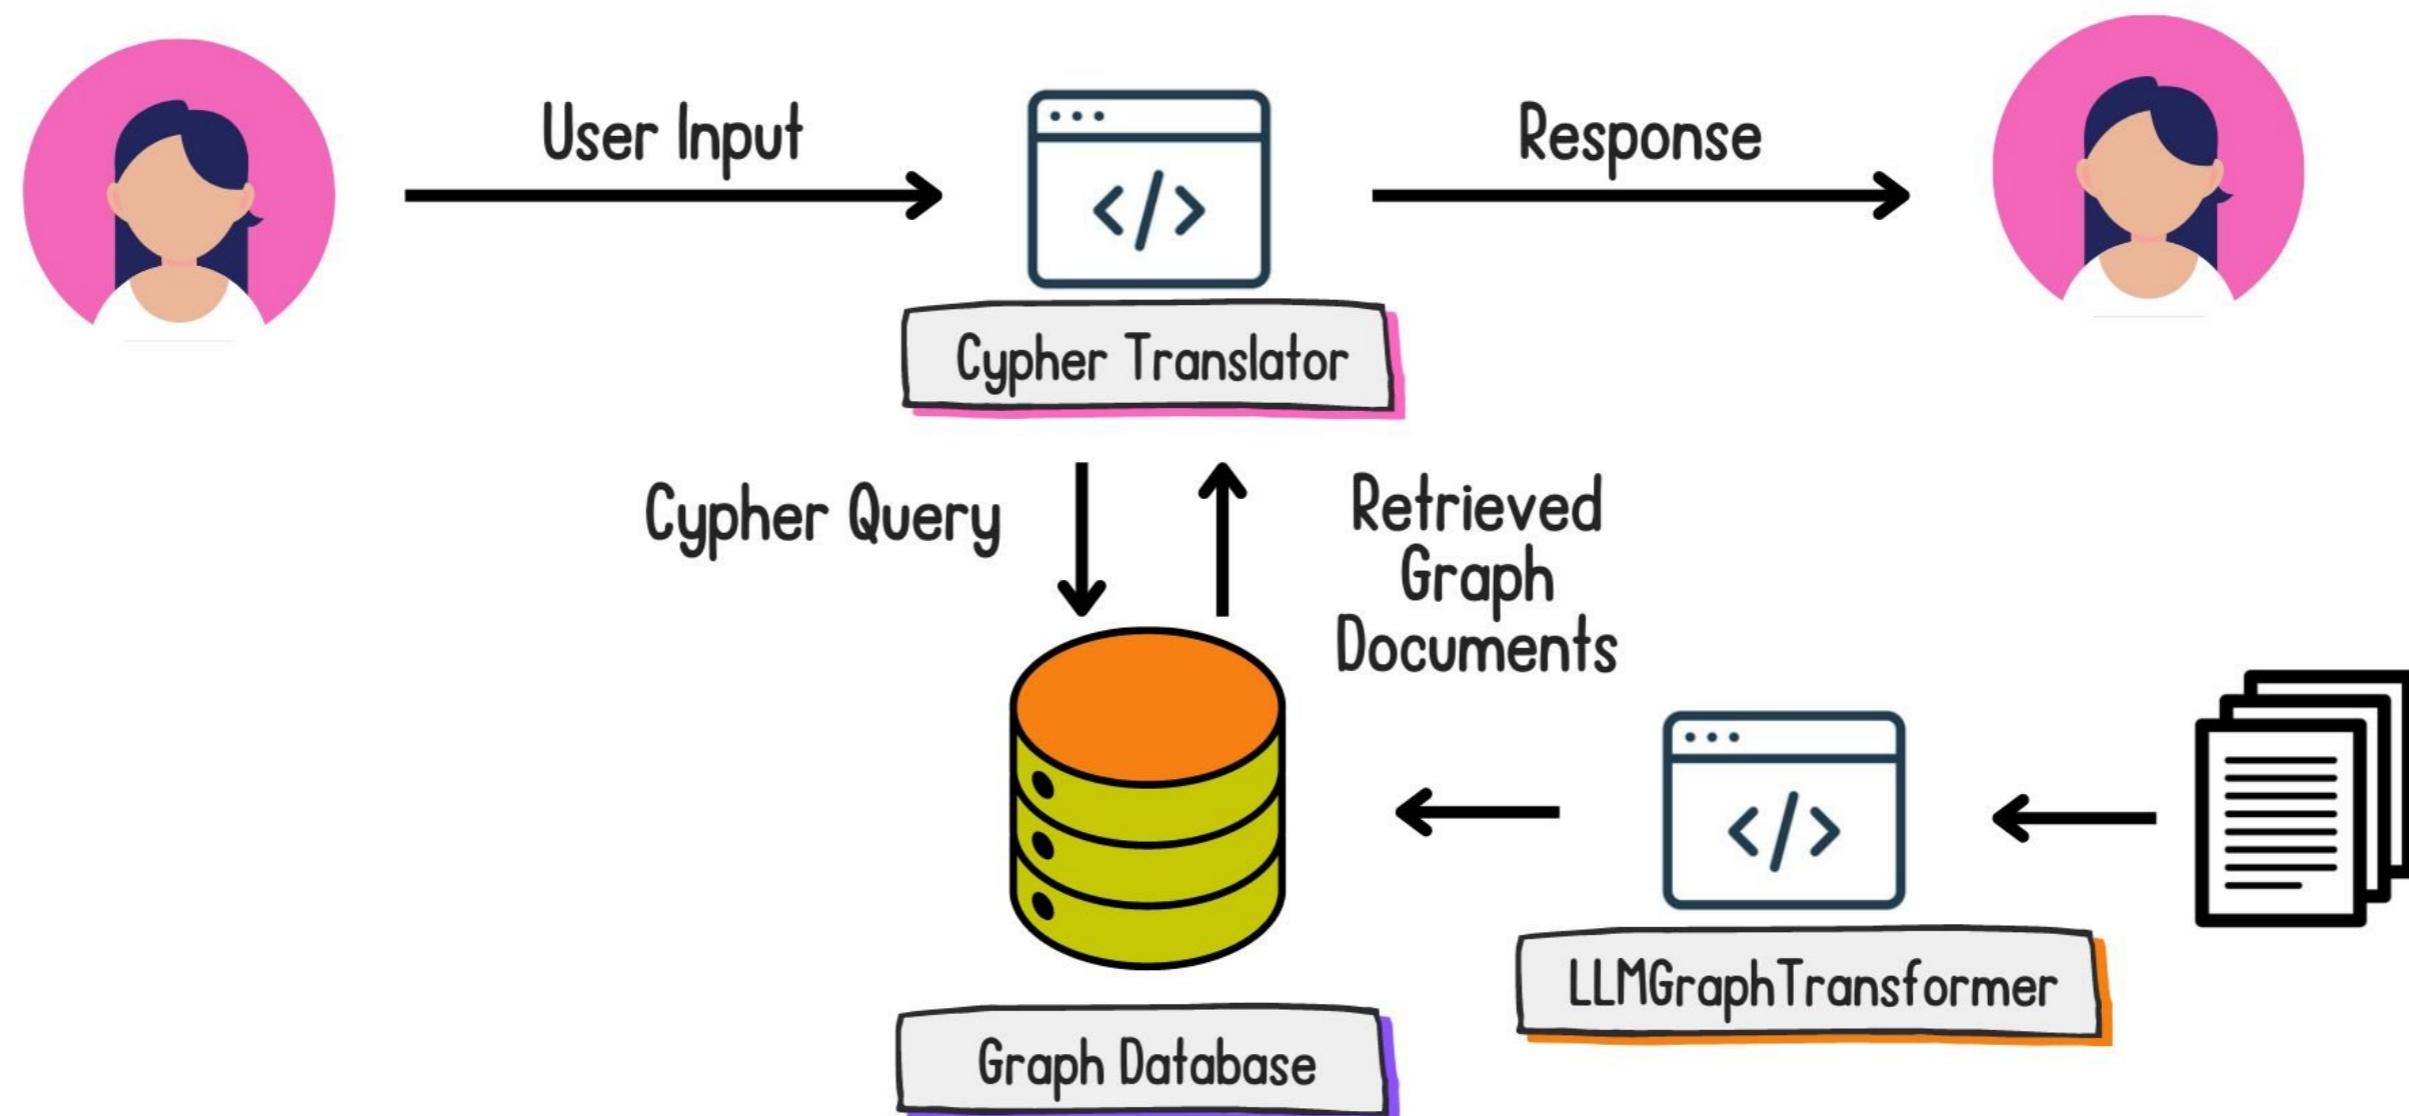
# Building the Graph RAG architecture



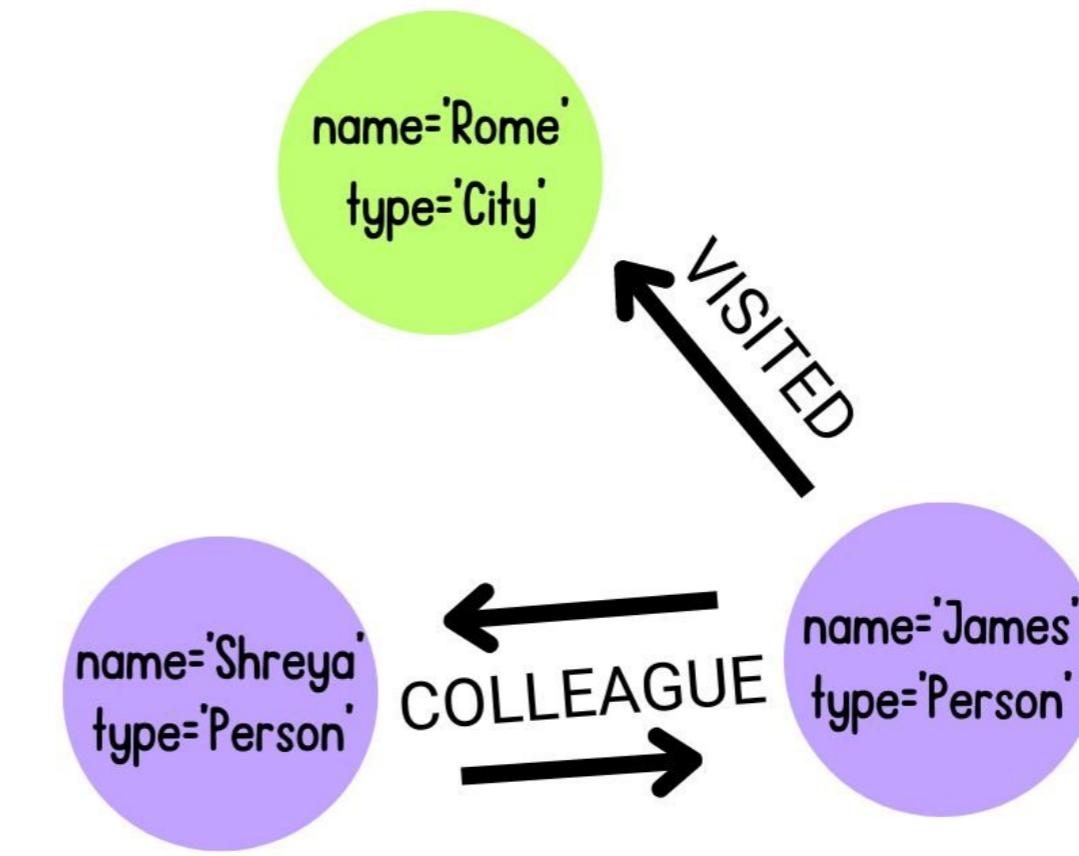
# Building the Graph RAG architecture



# Building the Graph RAG architecture



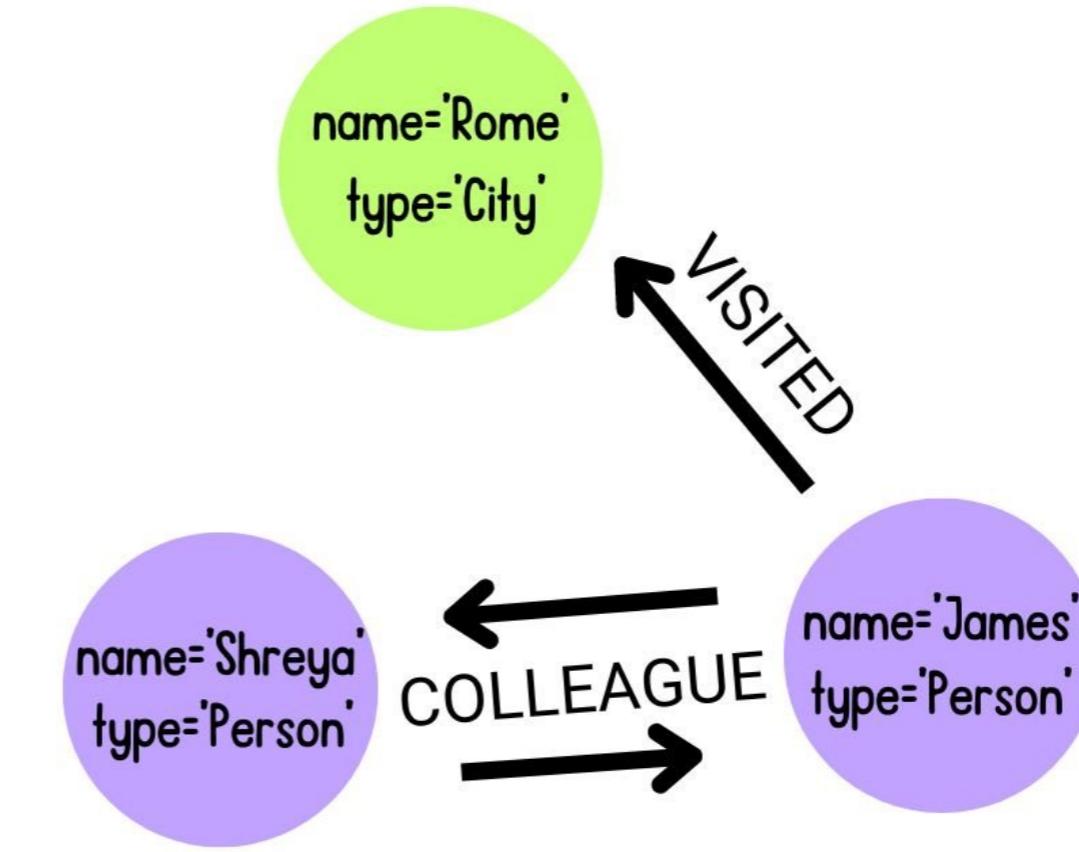
# From user inputs to Cypher queries



# From user inputs to Cypher queries

Where has James visited?

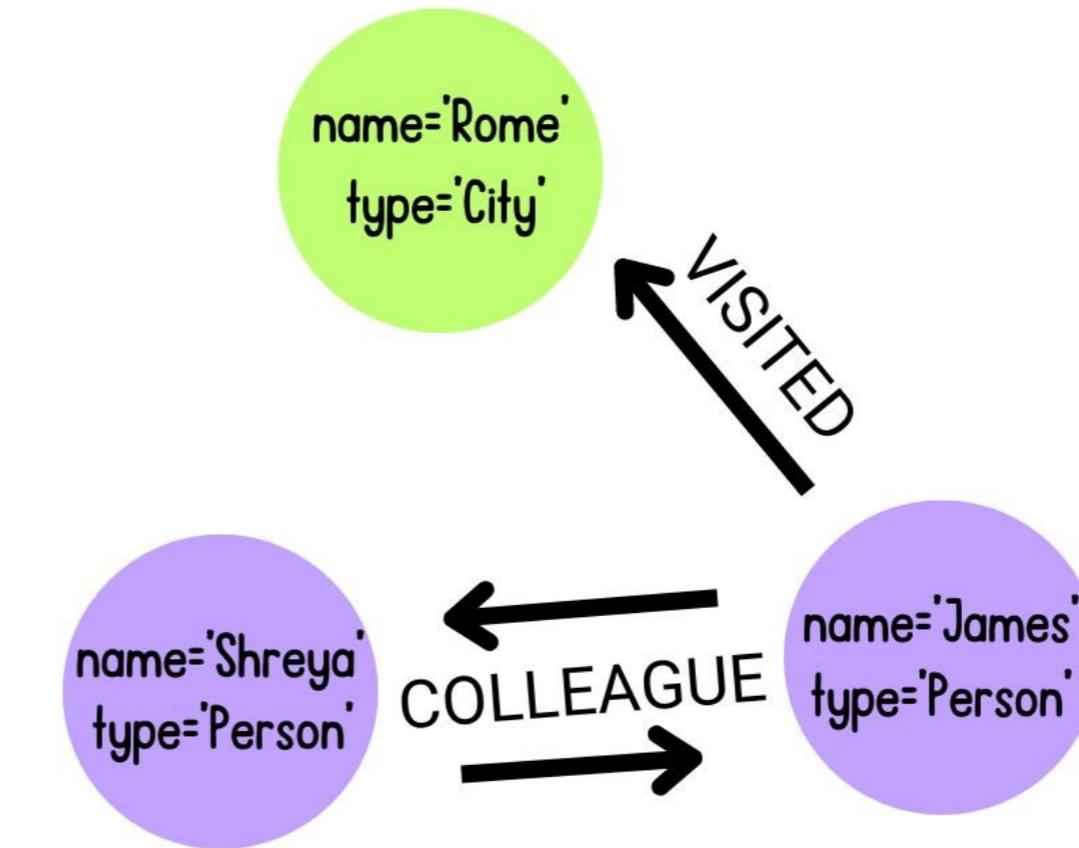
+



# From user inputs to Cypher queries

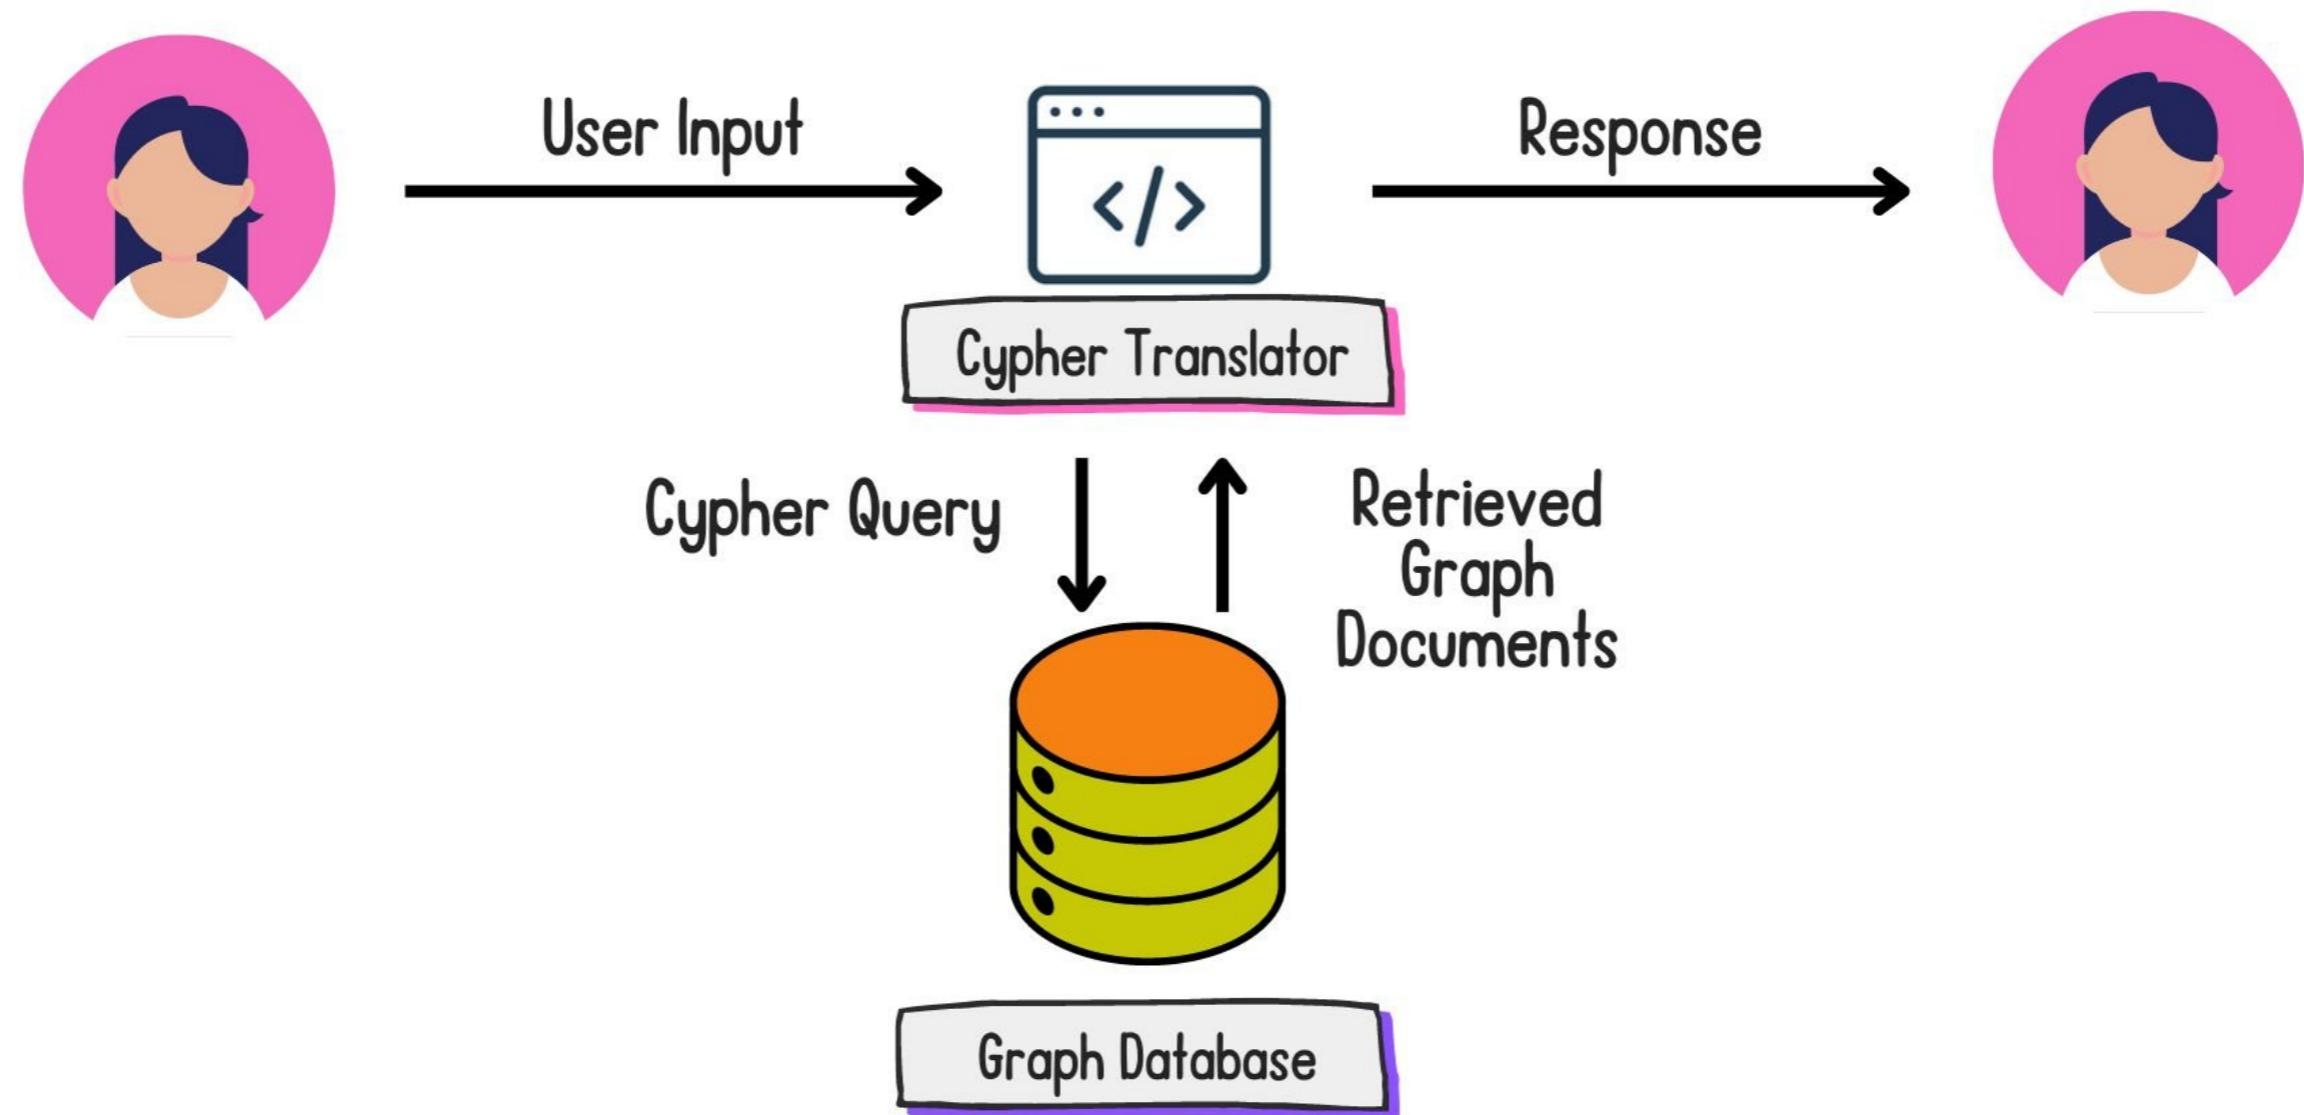
Where has James visited?

+

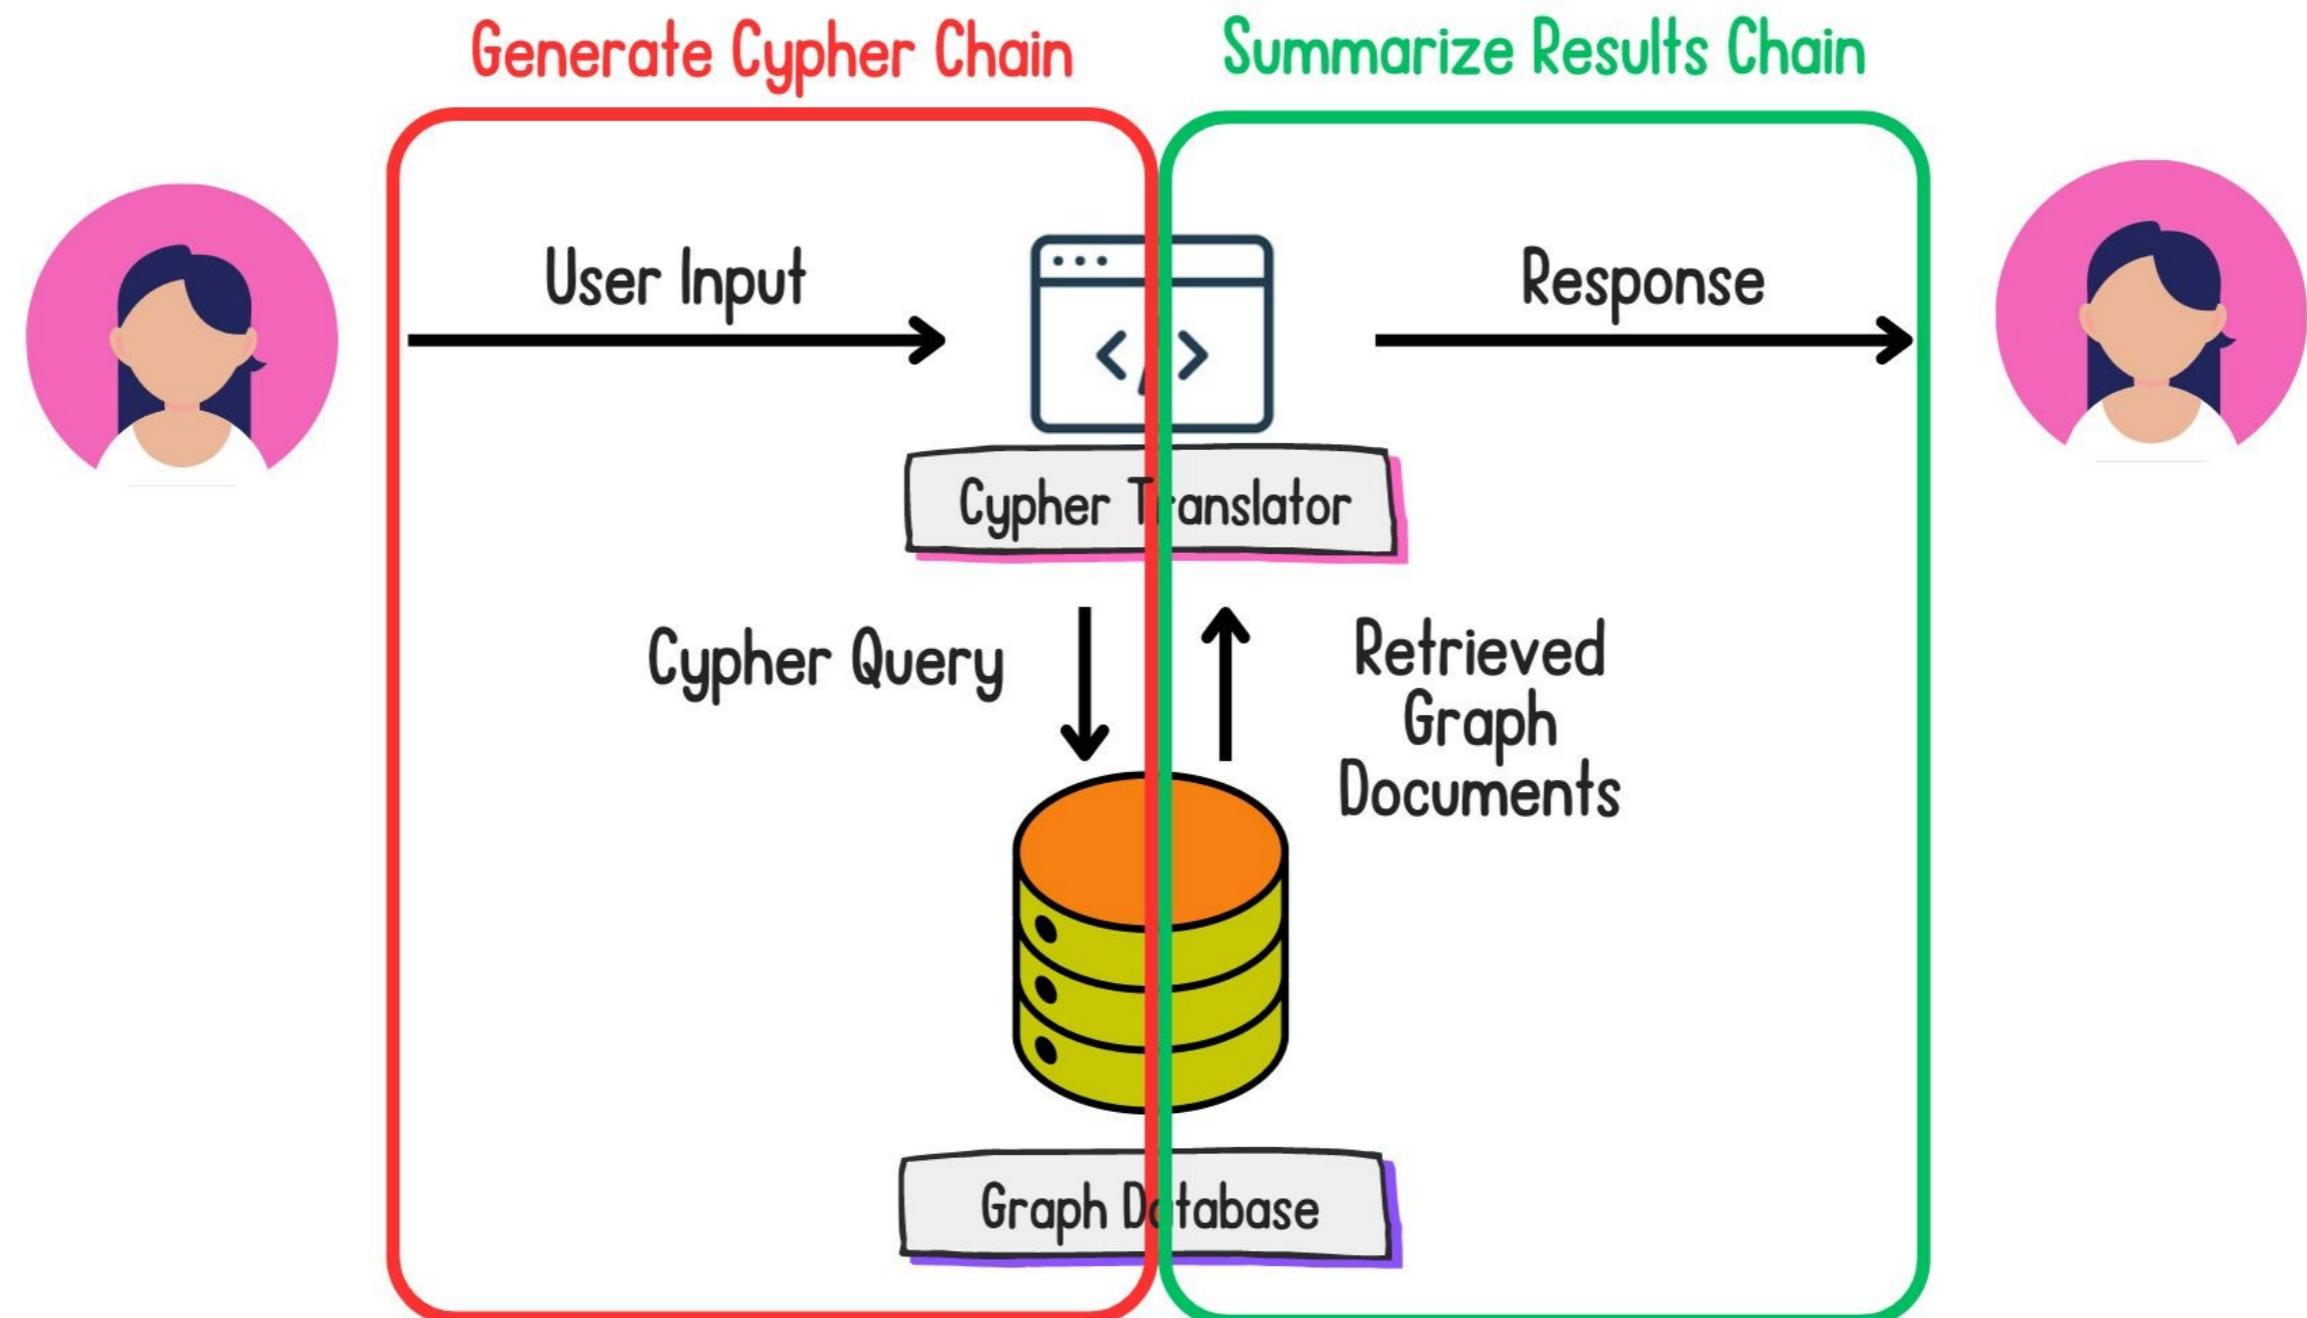


```
MATCH (james:Person{name:"James"}) -[:VISITED]> (location)
RETURN location
```

# GraphCypherQAChain



# GraphCypherQAChain



# Refresh schema

```
graph.refresh_schema()  
print(graph.get_schema)
```

Node properties:

Document {title: STRING, id: STRING, text: STRING, summary: STRING, source: STRING}

Concept {id: STRING}

Organization {id: STRING}

Relationship properties:

The relationships:

(:Document)-[:MENTIONS]->(:Organization)

(:Concept)-[:DEVELOPED\_BY]->(:Person)

# Querying the graph

```
from langchain_community.chains.graph_qa.cypher import GraphCypherQACChain

chain = GraphCypherQACChain.from_llm(
    llm=ChatOpenAI(api_key="...", temperature=0), graph=graph, verbose=True
)

result = chain.invoke({"query": "What is the most accurate model?"})
```

<sup>1</sup> [https://api.python.langchain.com/en/latest/chains/langchain\\_community.chains.graph\\_qa.cypher.GraphCypherQACChain.html](https://api.python.langchain.com/en/latest/chains/langchain_community.chains.graph_qa.cypher.GraphCypherQACChain.html)

# Querying the graph

```
print(f"Final answer: {result['result']}")
```

> Entering new GraphCypherQAChain chain...

Generated Cypher:

```
MATCH (m:Model)
RETURN m
ORDER BY m.accuracy DESC
LIMIT 1;
```

Full Context:

```
[{'m': {'id': 'Artificial Neural Networks'}}]
```

> Finished chain.

Final answer: Artificial Neural Networks

# Customization

```
chain = GraphCypherQACChain.from_llm(  
    llm=ChatOpenAI(api_key="...", temperature=0), graph=graph, verbose=True  
)
```

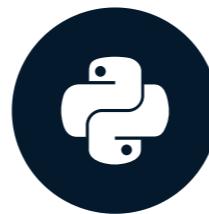
- `qa_prompt` : Prompt template for result generation
- `cypher_prompt` : Prompt template for Cypher generation
- `cypher_llm` : LLM for Cypher generation
- `qa_llm` : LLM for result generation

# **Let's practice!**

**RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN**

# Improving graph retrieval

RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN



**Meri Nova**  
Machine Learning Engineer

# Techniques

**Main limitation:** reliability of user → Cypher translation

Strategies to improve graph retrieval system:

- Filtering Graph Schema
- Validating the Cypher Query
- Few-shot prompting

# Filtering

```
from langchain_community.chains.graph_qa.cypher import GraphCypherQACChain

llm = ChatOpenAI(api_key="...", model="gpt-4o-mini", temperature=0)

chain = GraphCypherQACChain.from_llm(
    graph=graph, llm=llm, exclude_types=["Concept"], verbose=True
)
print(graph.get_schema)
```

Node properties:

Document {title: STRING, id: STRING, text: STRING, summary: STRING, source: STRING}

Organization {id: STRING}

# Validating the Cypher query

- Difficulty in interpreting the *direction* of relationships

```
chain = GraphCypherQAChain.from_llm(  
    graph=graph, llm=llm, verbose=True, validate_cypher=True  
)
```

1. Detects nodes and relationships
2. Determines the directions of the relationship
3. Checks the graph schema
4. Update the direction of relationships

# Few-shot prompting

```
examples = [
    {
        "question": "How many notable large language models are mentioned in the article?",
        "query": "MATCH (m:Concept {id: 'Large Language Model'}) RETURN count(DISTINCT m)",
    },
    {
        "question": "Which companies or organizations have developed the large language models mentioned?",
        "query": "MATCH (o:Organization)-[:DEVELOPS]->(m:Concept {id: 'Large Language Model'}) RETURN DISTINCT o.id",
    },
    {
        "question": "What is the largest model size mentioned in the article, in terms of number of parameters?",
        "query": "MATCH (m:Concept {id: 'Large Language Model'}) RETURN max(m.parameters) AS largest_model",
    },
]
```

# Implementing few-shot prompting

```
from langchain_core.prompts import FewShotPromptTemplate, PromptTemplate

example_prompt = PromptTemplate.from_template(
    "User input: {question}\nCypher query: {query}"
)

cypher_prompt = FewShotPromptTemplate(
    examples=examples,
    example_prompt=example_prompt,
    prefix="You are a Neo4j expert. Given an input question, create a syntactically correct Cypher query to achieve the goal. Don't include the 'Cypher query:' prefix in your response.",
    suffix="User input: {question}\nCypher query: ",
    input_variables=["question"],
)
```

# Complete prompt

You are a Neo4j expert. Given an input question, create a syntactically correct Cypher query to run.

Below are a number of examples of questions and their corresponding Cypher queries.

User input: How many notable large language models are mentioned in the article?

Cypher query: MATCH (p:Paper) RETURN count(DISTINCT p)

User input: Which companies or organizations have developed the large language models?

Cypher query: MATCH (o:Organization)-[:DEVELOPS]->(m:Concept {id: 'Large Language Model'}) RETURN DISTINCT o.id

User input: What is the largest model size mentioned in the article, in terms of number of parameters?

Cypher query: MATCH (m:Concept {id: 'Large Language Model'}) RETURN max(m.parameters) AS largest\_model

User input: How many papers were published in 2016?

Cypher query:

# Adding the few-shot examples

```
chain = GraphCypherQACChain.from_llm(  
    graph=graph, llm=llm, cypher_prompt=cypher_prompt,  
    verbose=True, validate_cypher=True  
)
```

# **Let's practice!**

**RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN**

# Congratulations!

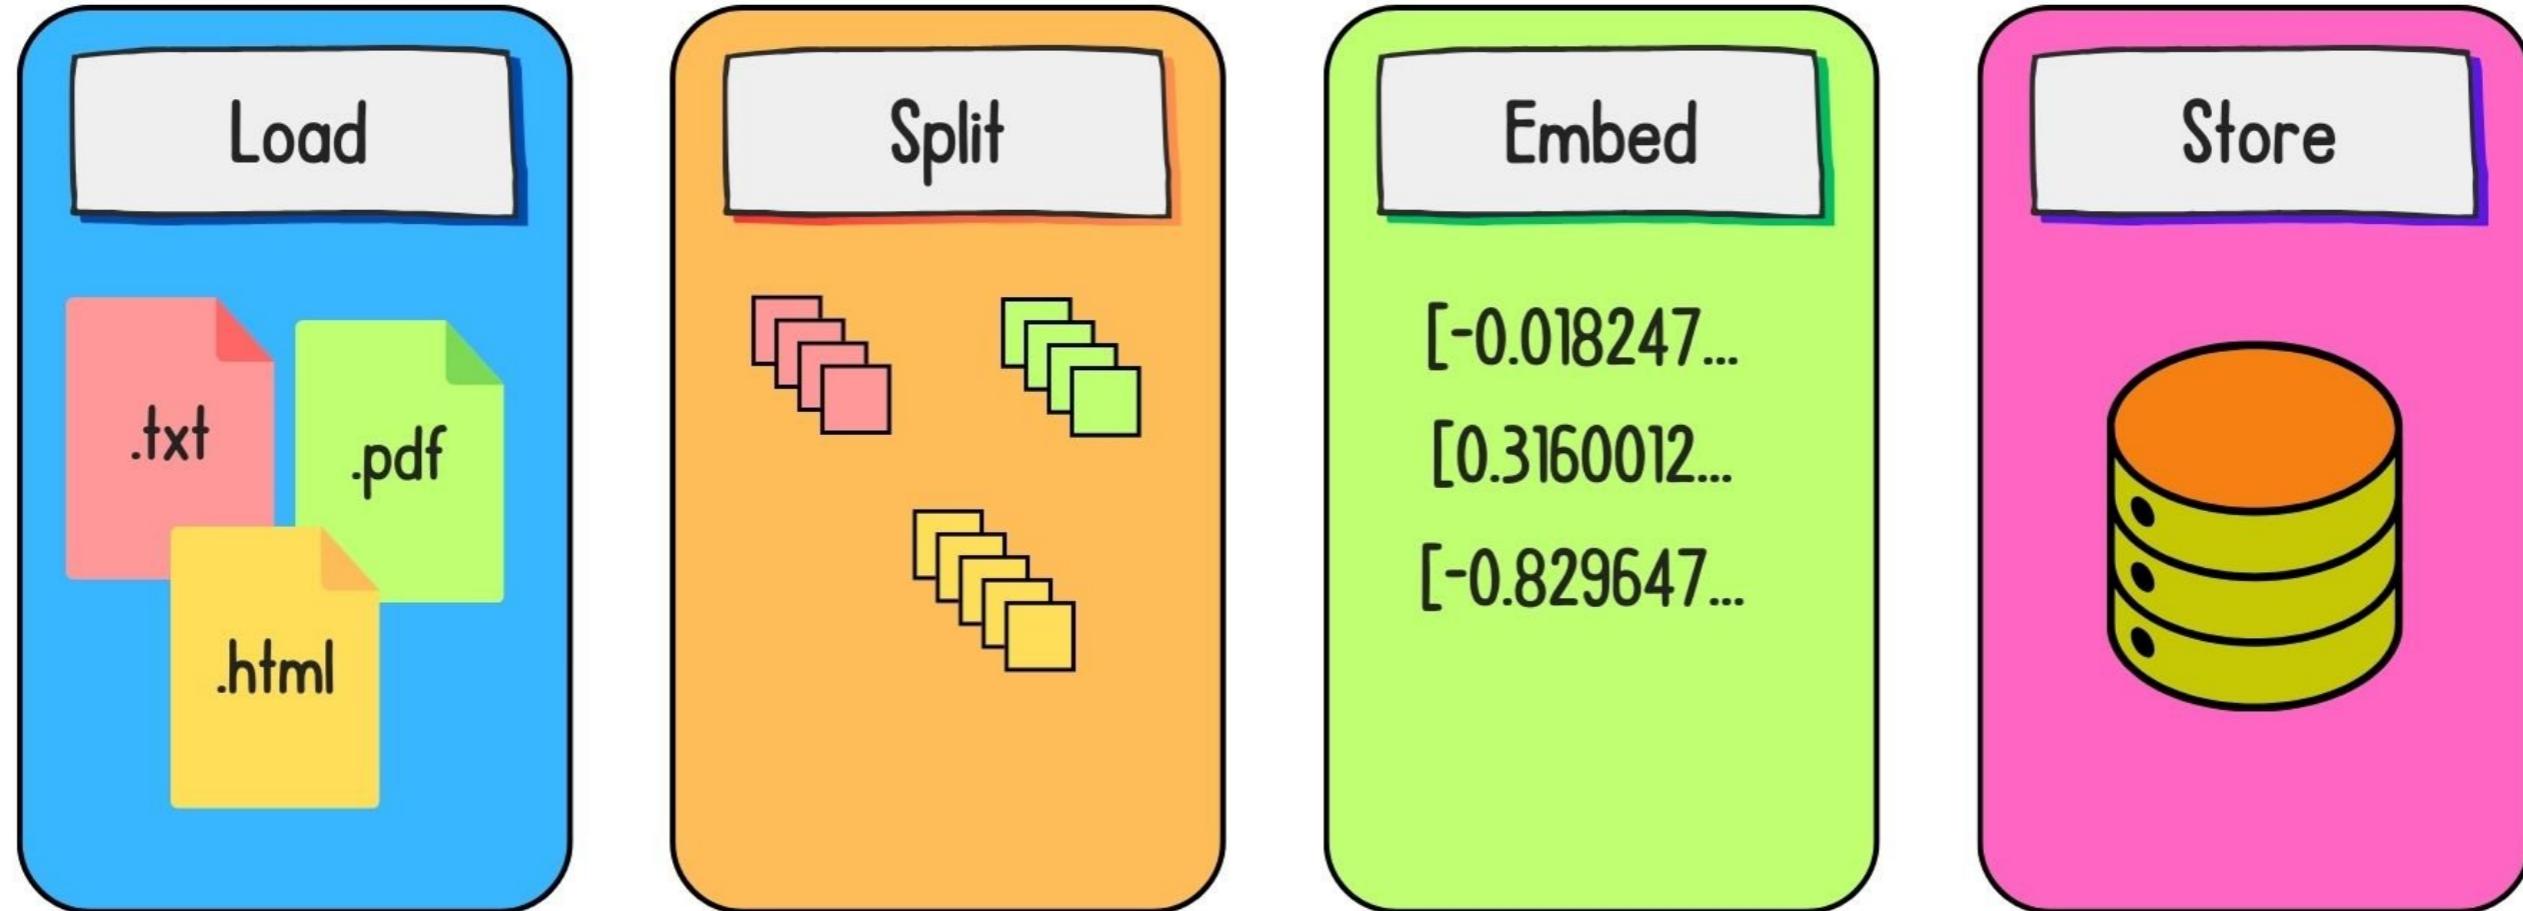
RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN



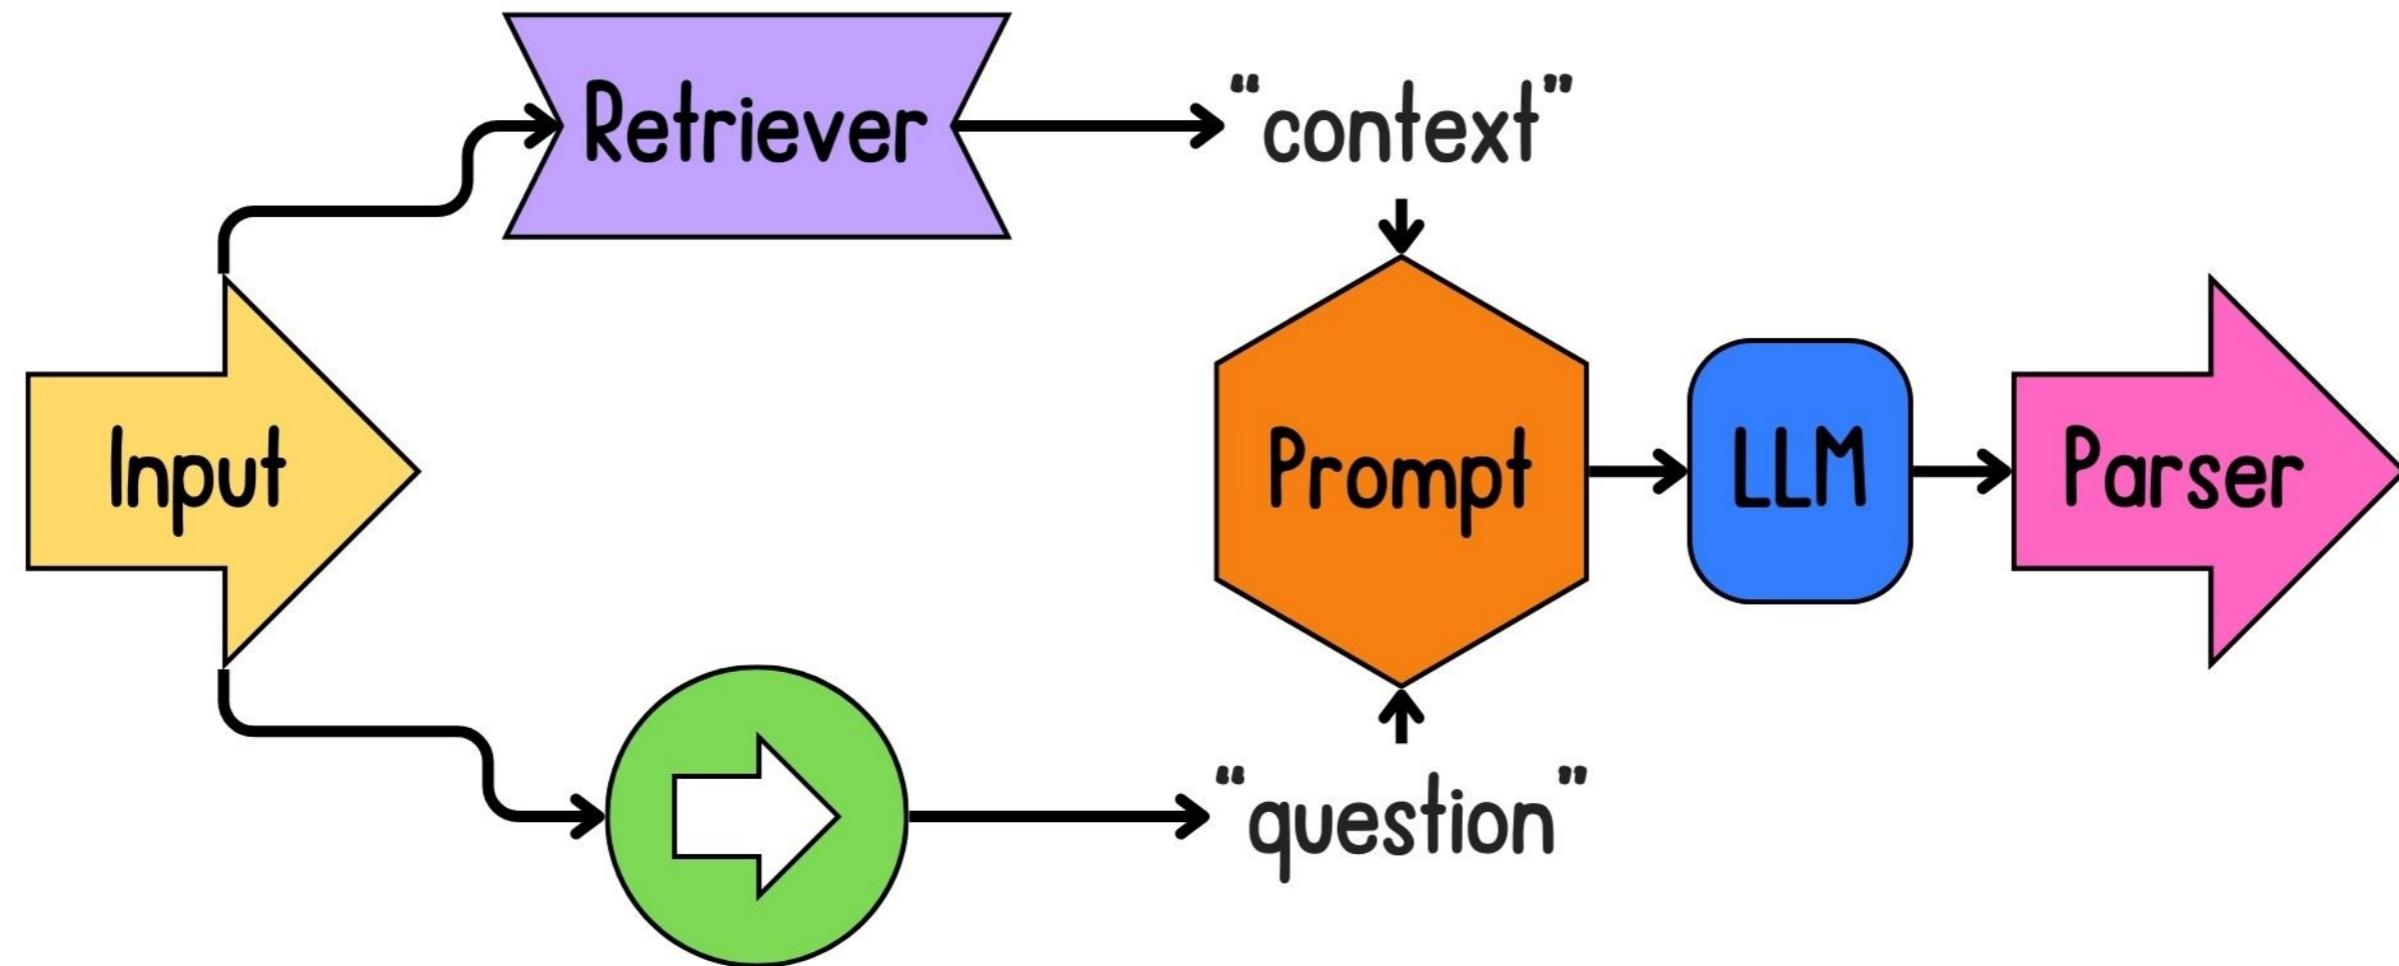
Meri Nova

Machine Learning Engineer

# Chapter 1



# Chapter 1



RunnablePassthrough

# Chapter 2

## Markdown files

- UnstructuredMarkdownLoader

## Python files

- PythonLoader
- language=Language.PYTHON

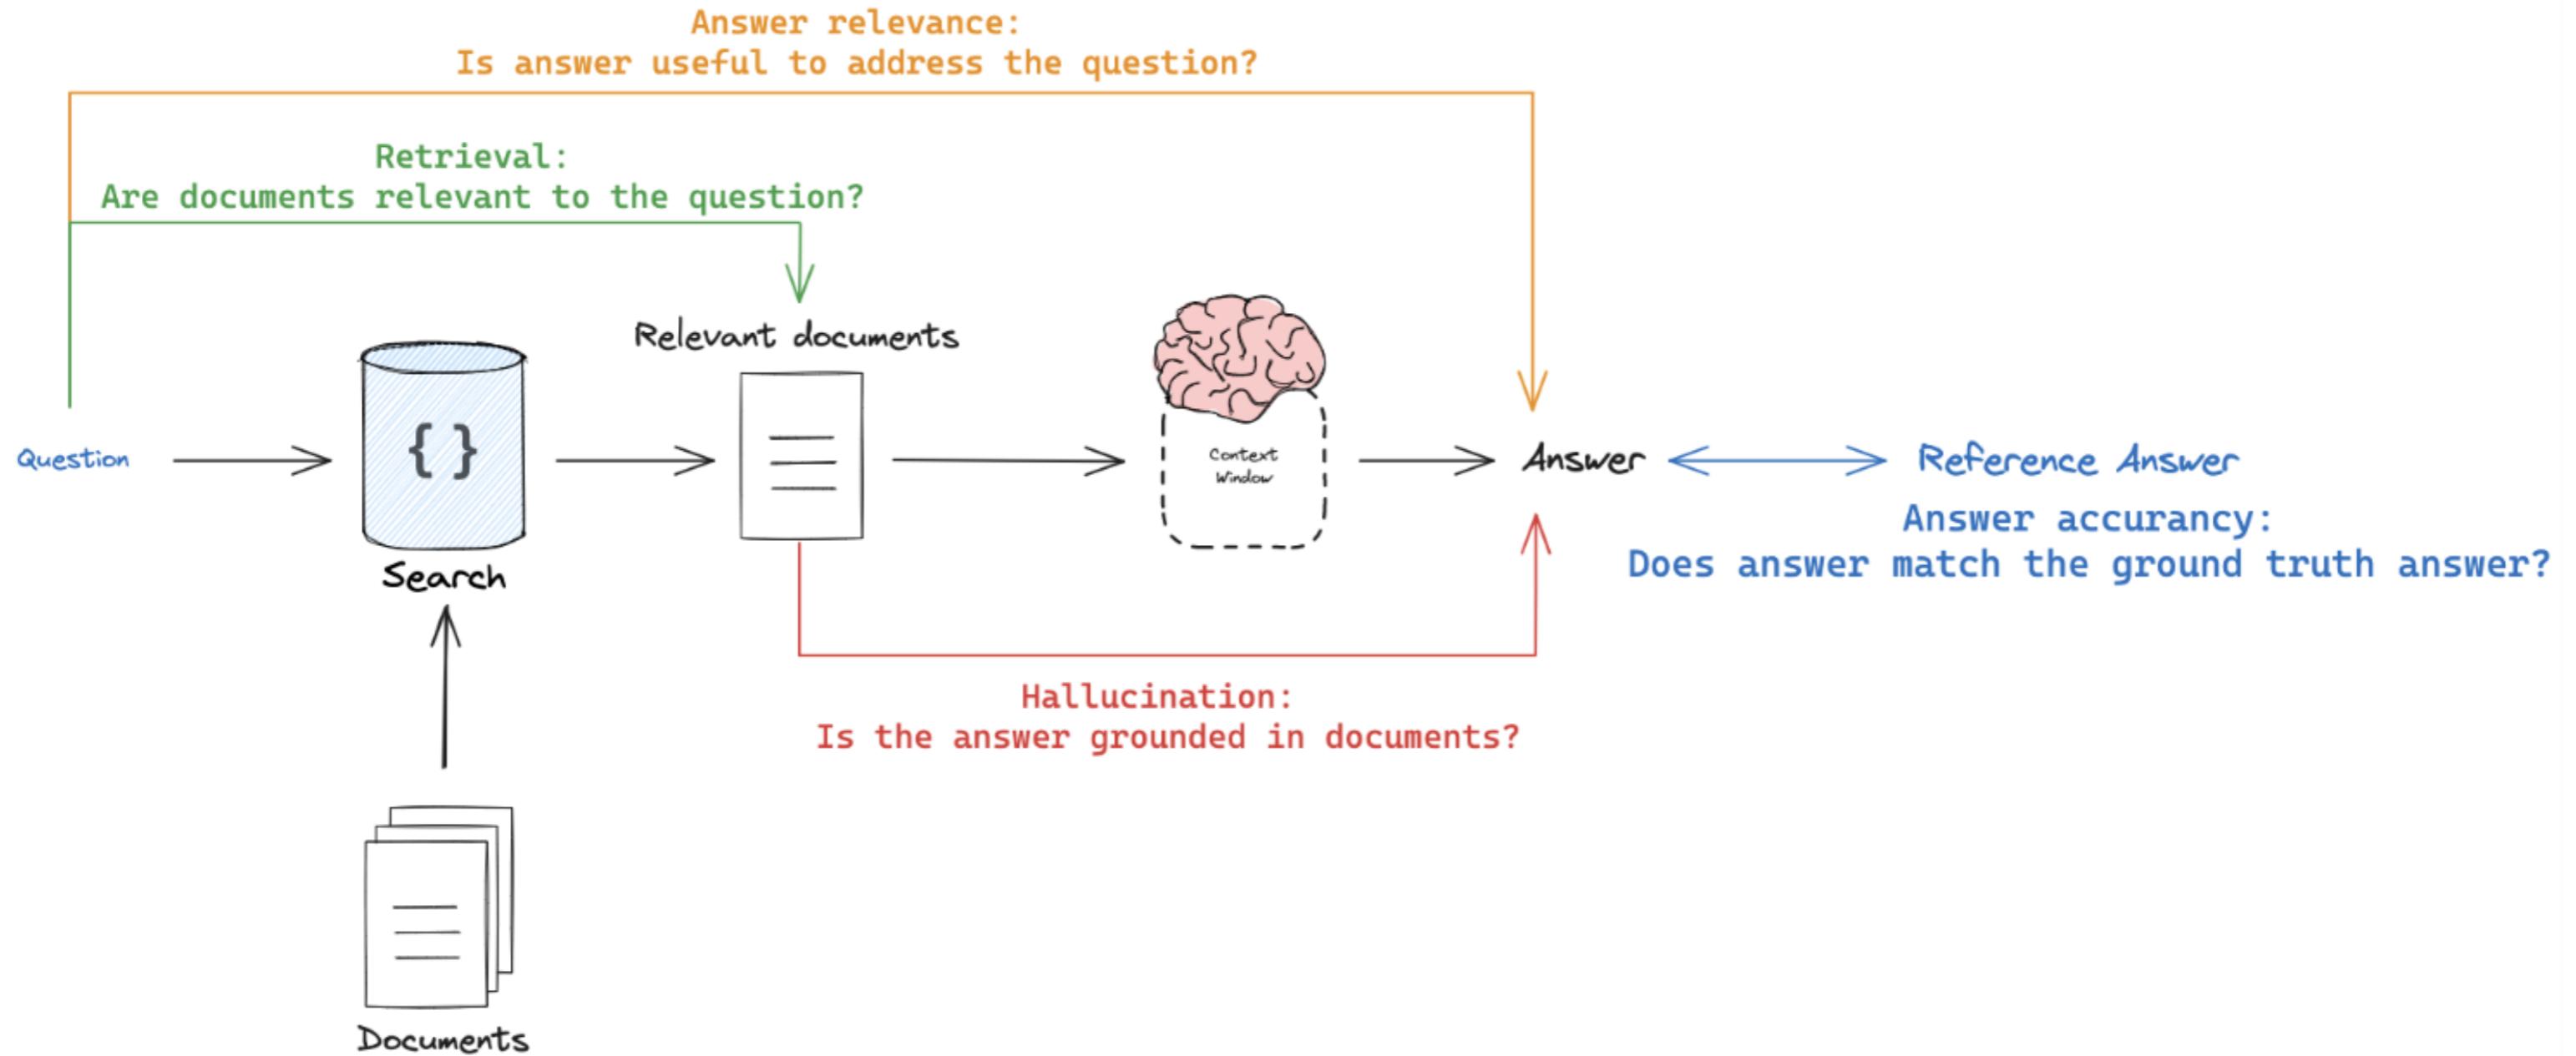
Token splitting → TokenTextSplitter

Semantic splitting → SemanticChunker

RAG applications has numerous use cases,  
but are most frequently deployed in chatbots.

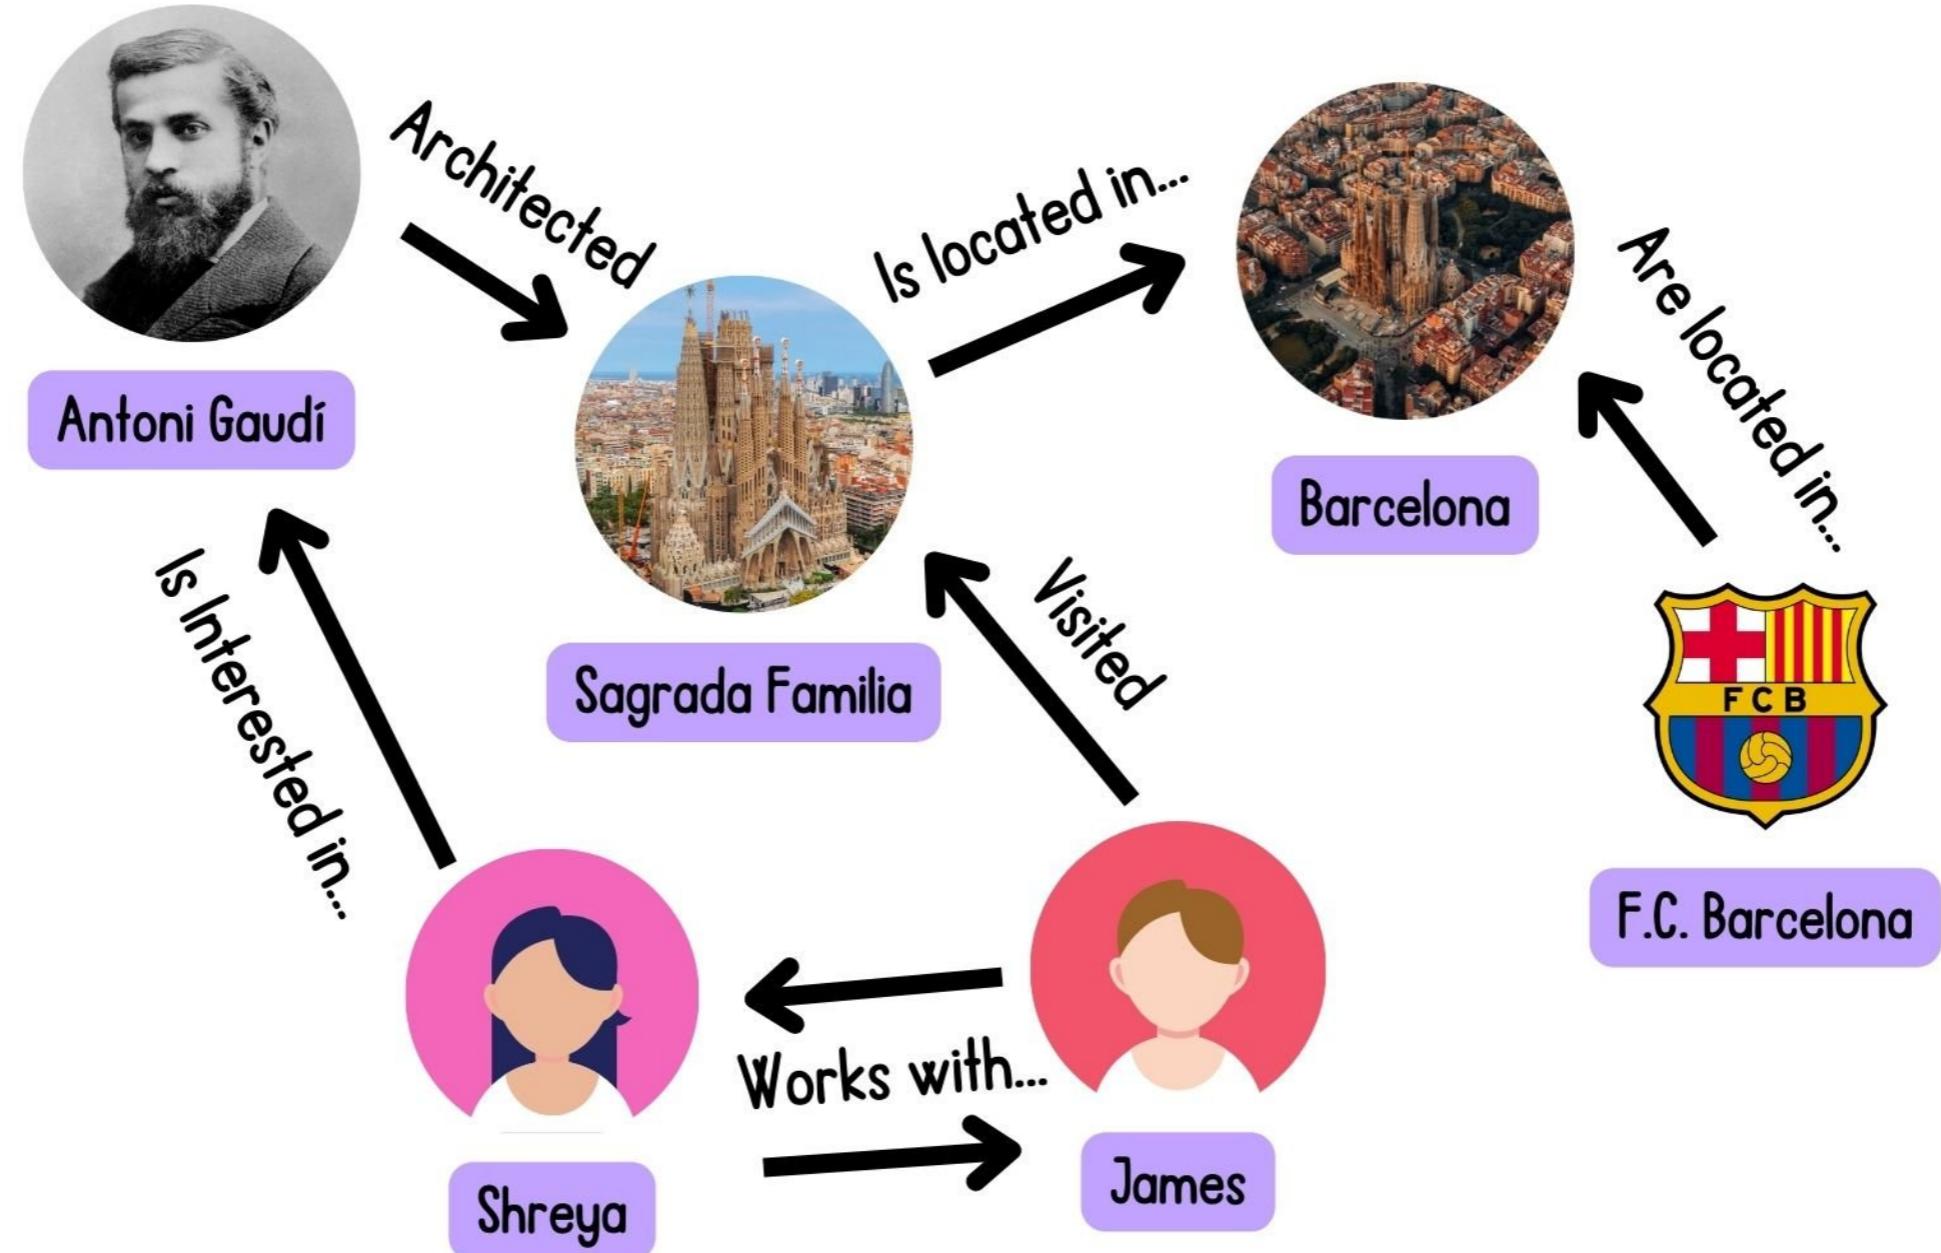
Domestic dogs are descendants from an  
extinct population of Pleistocene wolves over  
14,000 years ago.

# Chapter 2

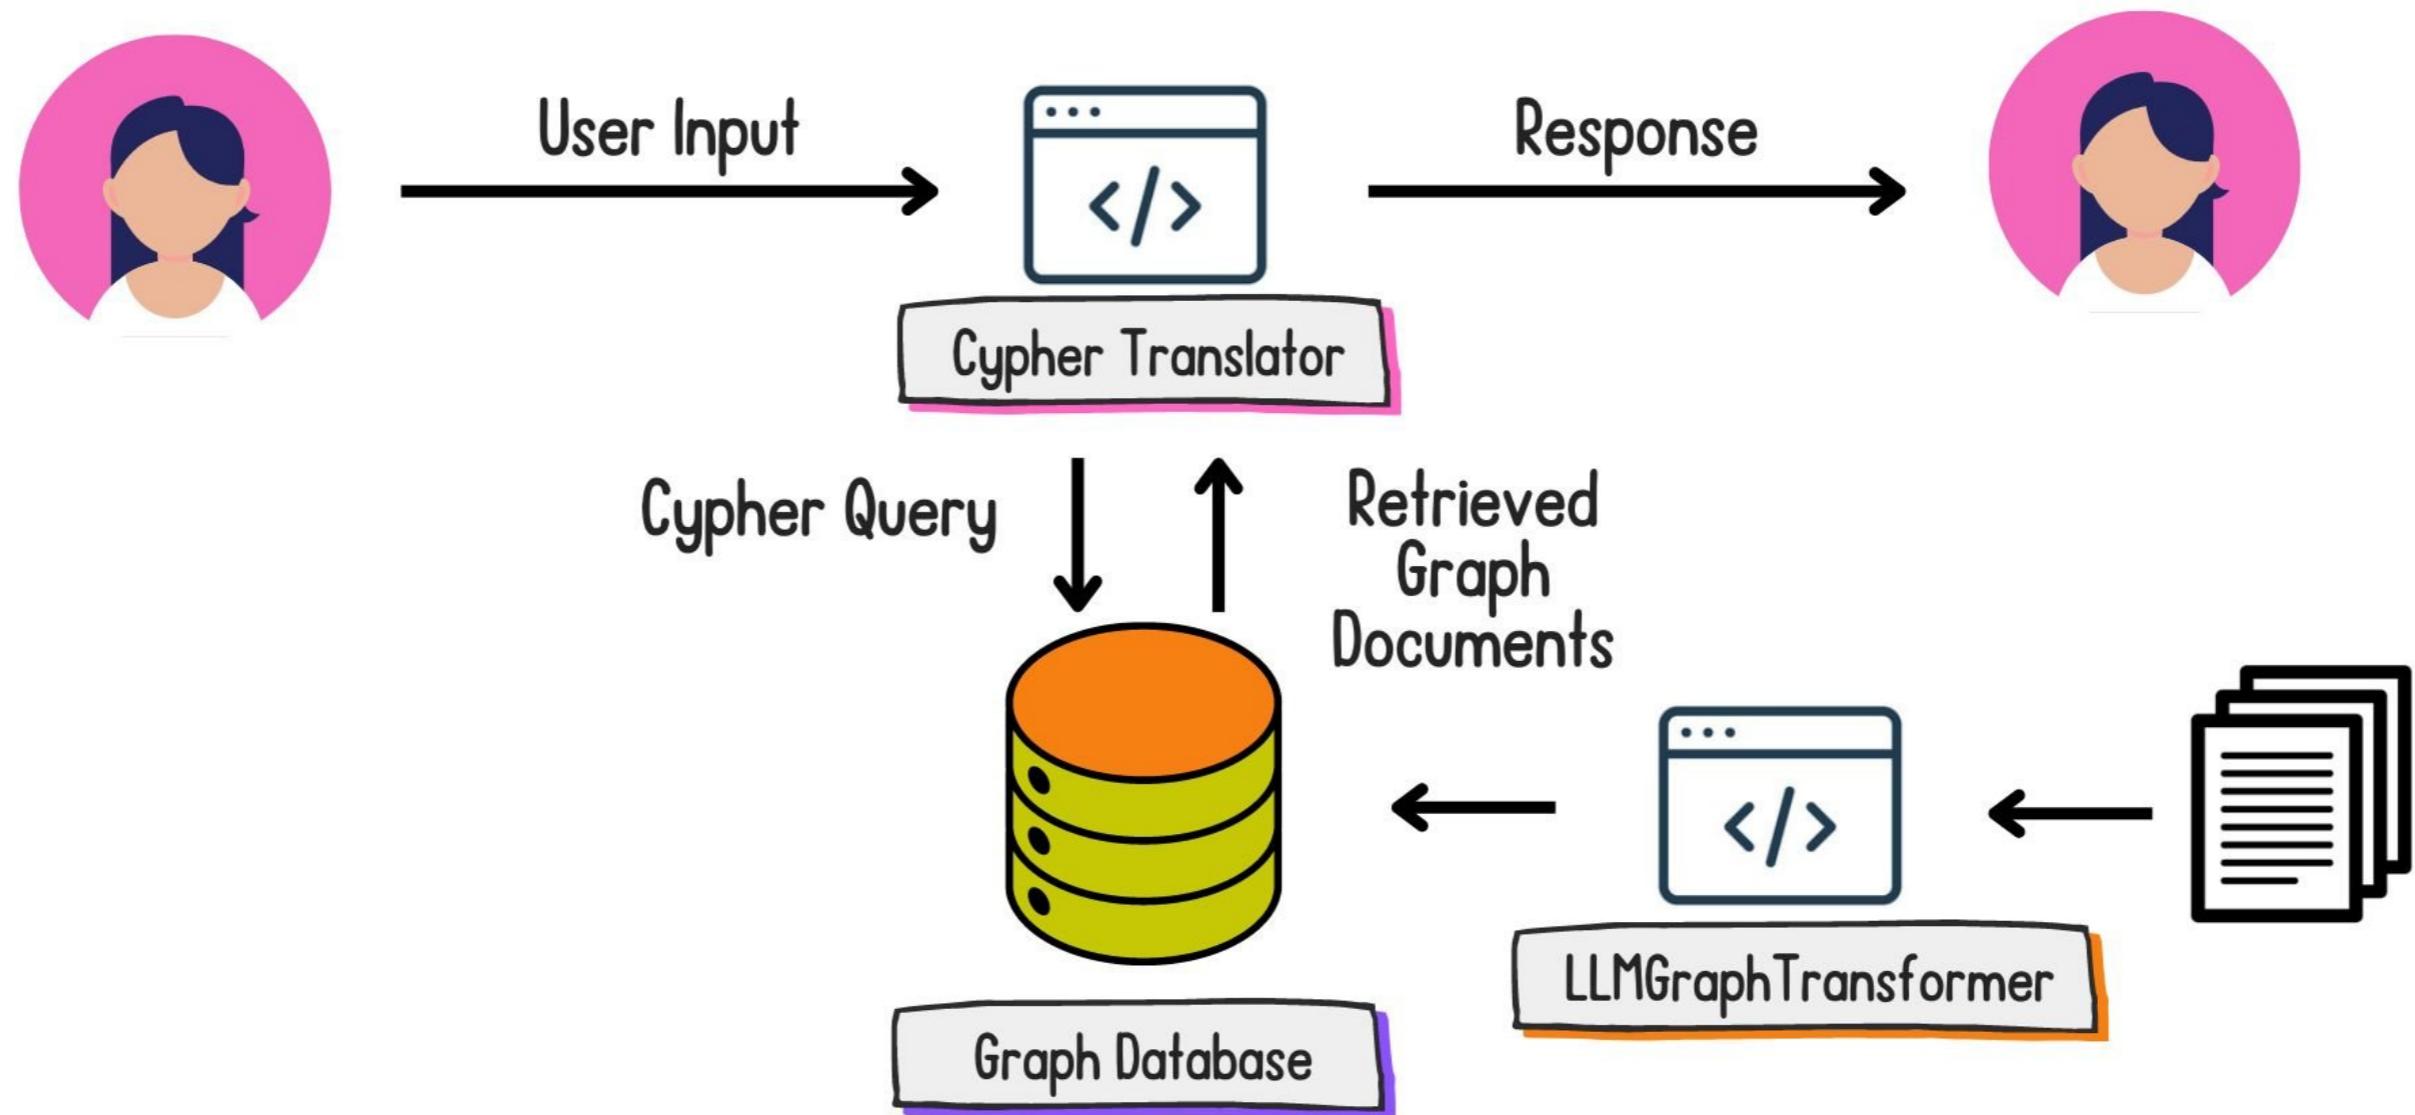


<sup>1</sup> Image Credit: LangSmith

# Chapter 3



# Chapter 3



# **Let's practice!**

**RETRIEVAL AUGMENTED GENERATION (RAG) WITH LANGCHAIN**